

DATASTAX:

Streaming Big Data with Spark, Kafka,
Cassandra, Akka & Scala

Helena Edelson
[@helenaedelson](https://twitter.com/helenaedelson)



1

Delivering Meaning In Near-Real Time At High Velocity

2

Overview Of Spark Streaming, Kafka and Akka

3

Cassandra and the Spark Cassandra Connector

4

Integration In Big Data Applications

Who Is This Person?

- Spark Cassandra Connector Committer
- Akka Contributor (2 features in Cluster)
- Scala & Big Data Conference Speaker
-  @helenaedelson
-  <https://github.com/helena>
- Senior Software Engineer, Analytics @ DataStax
- Previously Senior Cloud Engineer at VMware & others

Use Case: Hadoop + Scalding

```
/** Reads SequenceFile data from S3 buckets, computes then persists to Cassandra. */
class TopSearches(args: Args) extends TopKDailyJob[MyDataType](args) with Cassandra {

  PailSource.source[Search](rootpath, structure, directories).read
    .mapTo('pailItem -> 'engines) { e: Search => results(e) }
    .filter('engines) { e: String => e.nonEmpty }
    .groupBy('engines) { _.size('count).sortBy('engines) }
    .groupBy('engines) { _.sortedReverseTake[(String, Long)]((('engines, 'count) -> 'tcount, k) }
    .flatMapTo('tcount -> ('key, 'engine, 'topCount)) { t: List[(String, Long)] =>
      t map { case (k, v) => (jobKey, k, v) }
    }
    .write(CassandraSource(connection, "top_searches", Scheme('key, ('engine, 'topCount))))}

}
```

Use Case: Spark?

```
20  object SparkWordCount extends WordCountBlueprint {  
21  
22      sc.textFile("./src/main/resources/data/words")  
23          .flatMap(_.split("\\s+"))  
24          .map(word => (clean(word), 1))  
25          .reduceByKey(_ + _)  
26          .saveAsTextFile(s"./wordcount-$timestamp")  
27  }
```

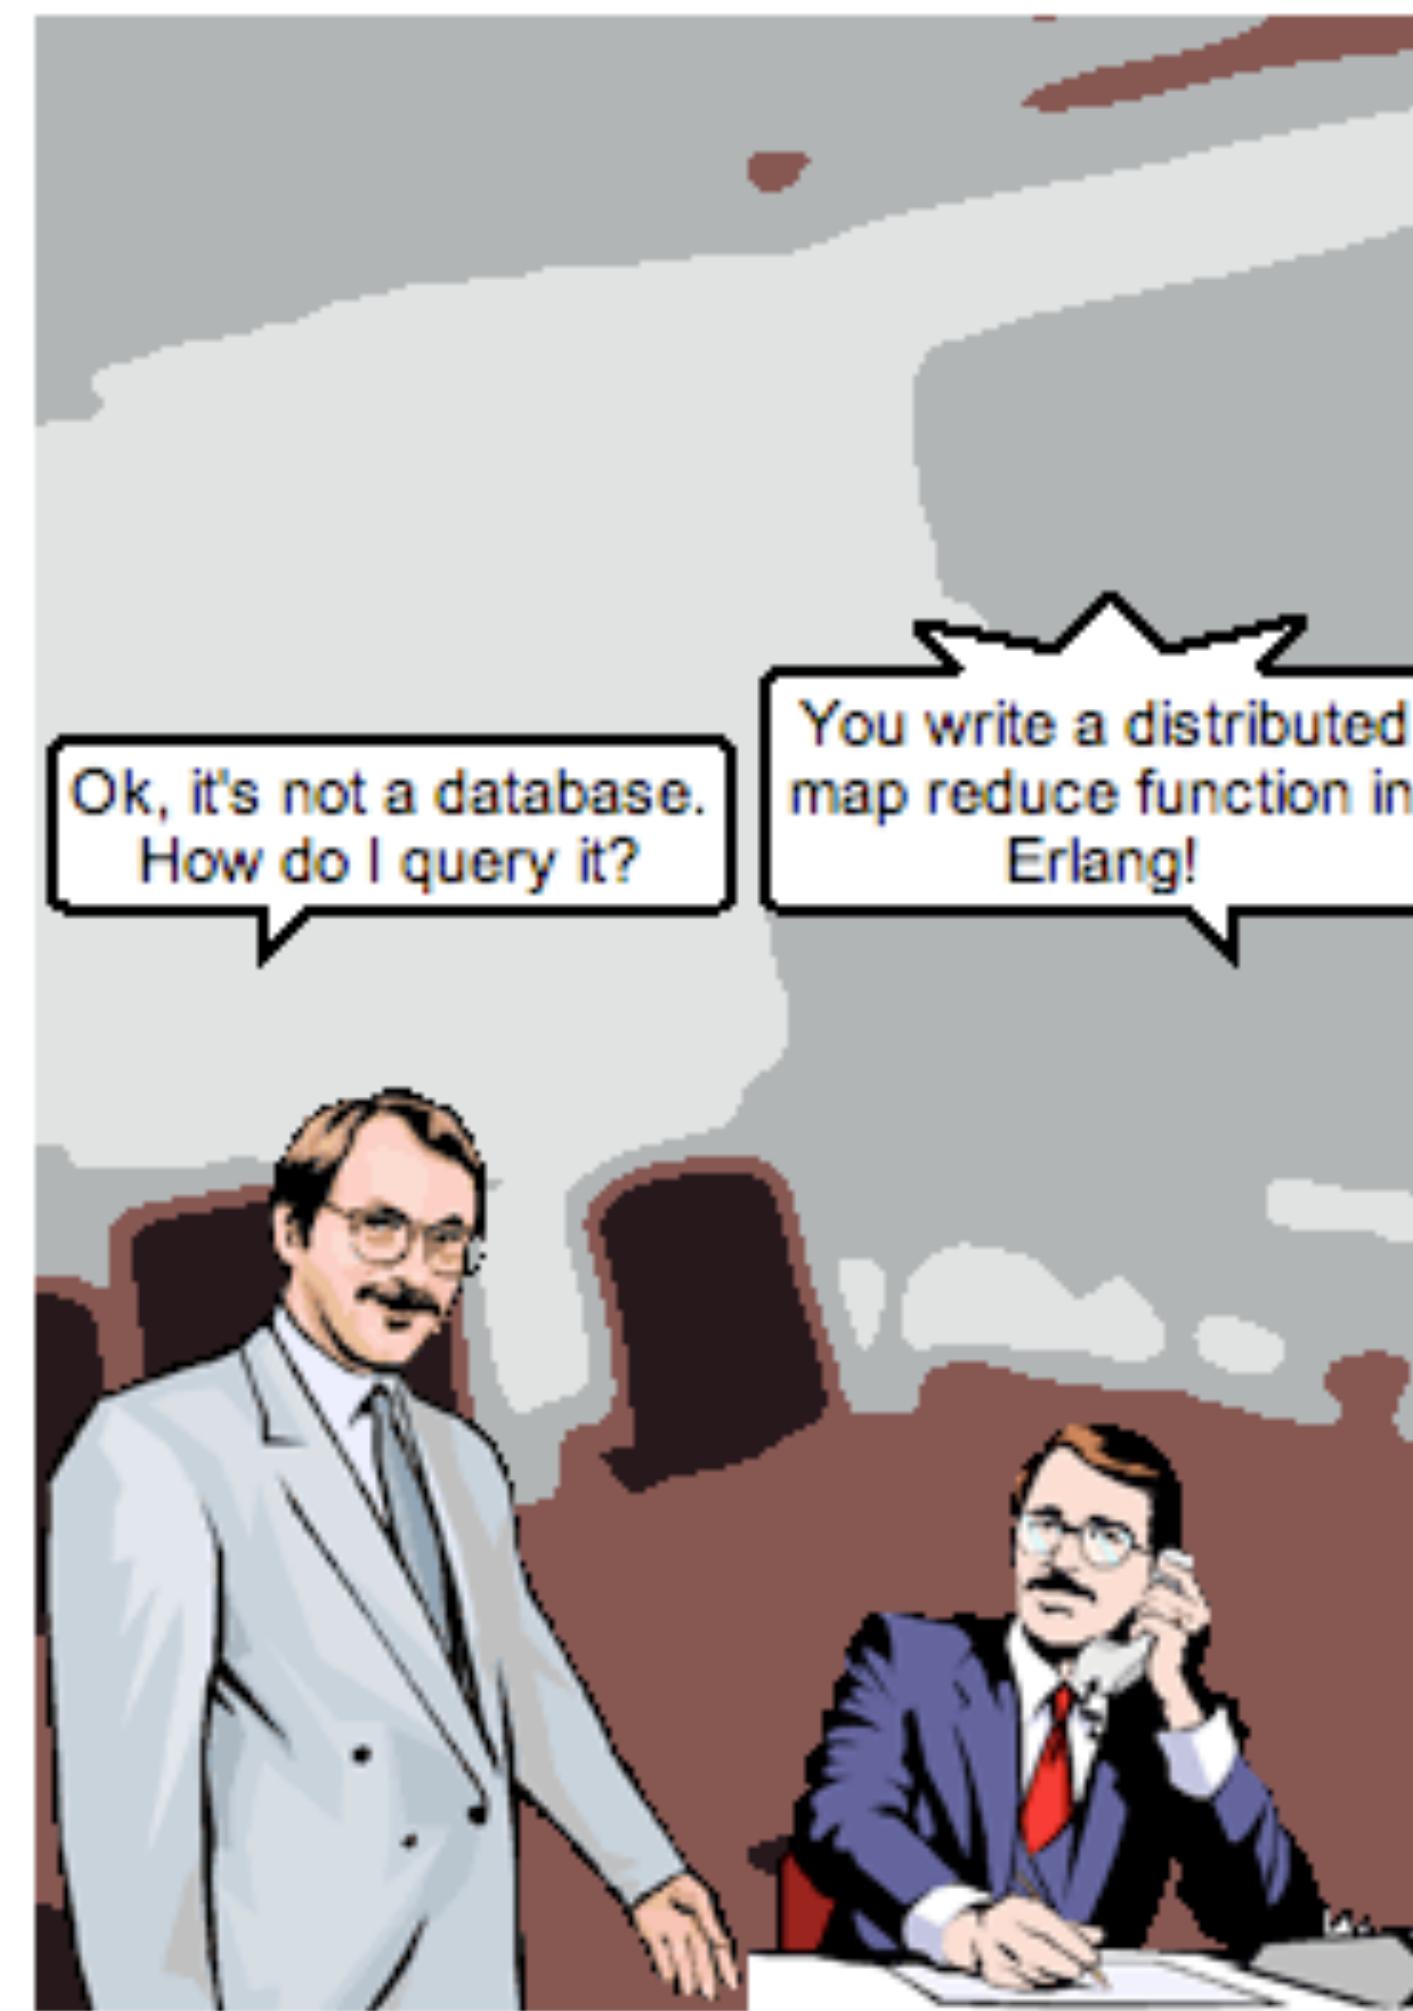
Delivering Meaning In Near-Real
Time At High Velocity At Massive
Scale

DATASTAX

Spark

Fault-tolerance

by @jrecursive



Strategies

- Partition For Scale
- Replicate For Resiliency
- Share Nothing
- Asynchronous Message Passing
- Parallelism
- Isolation
- Location Transparency

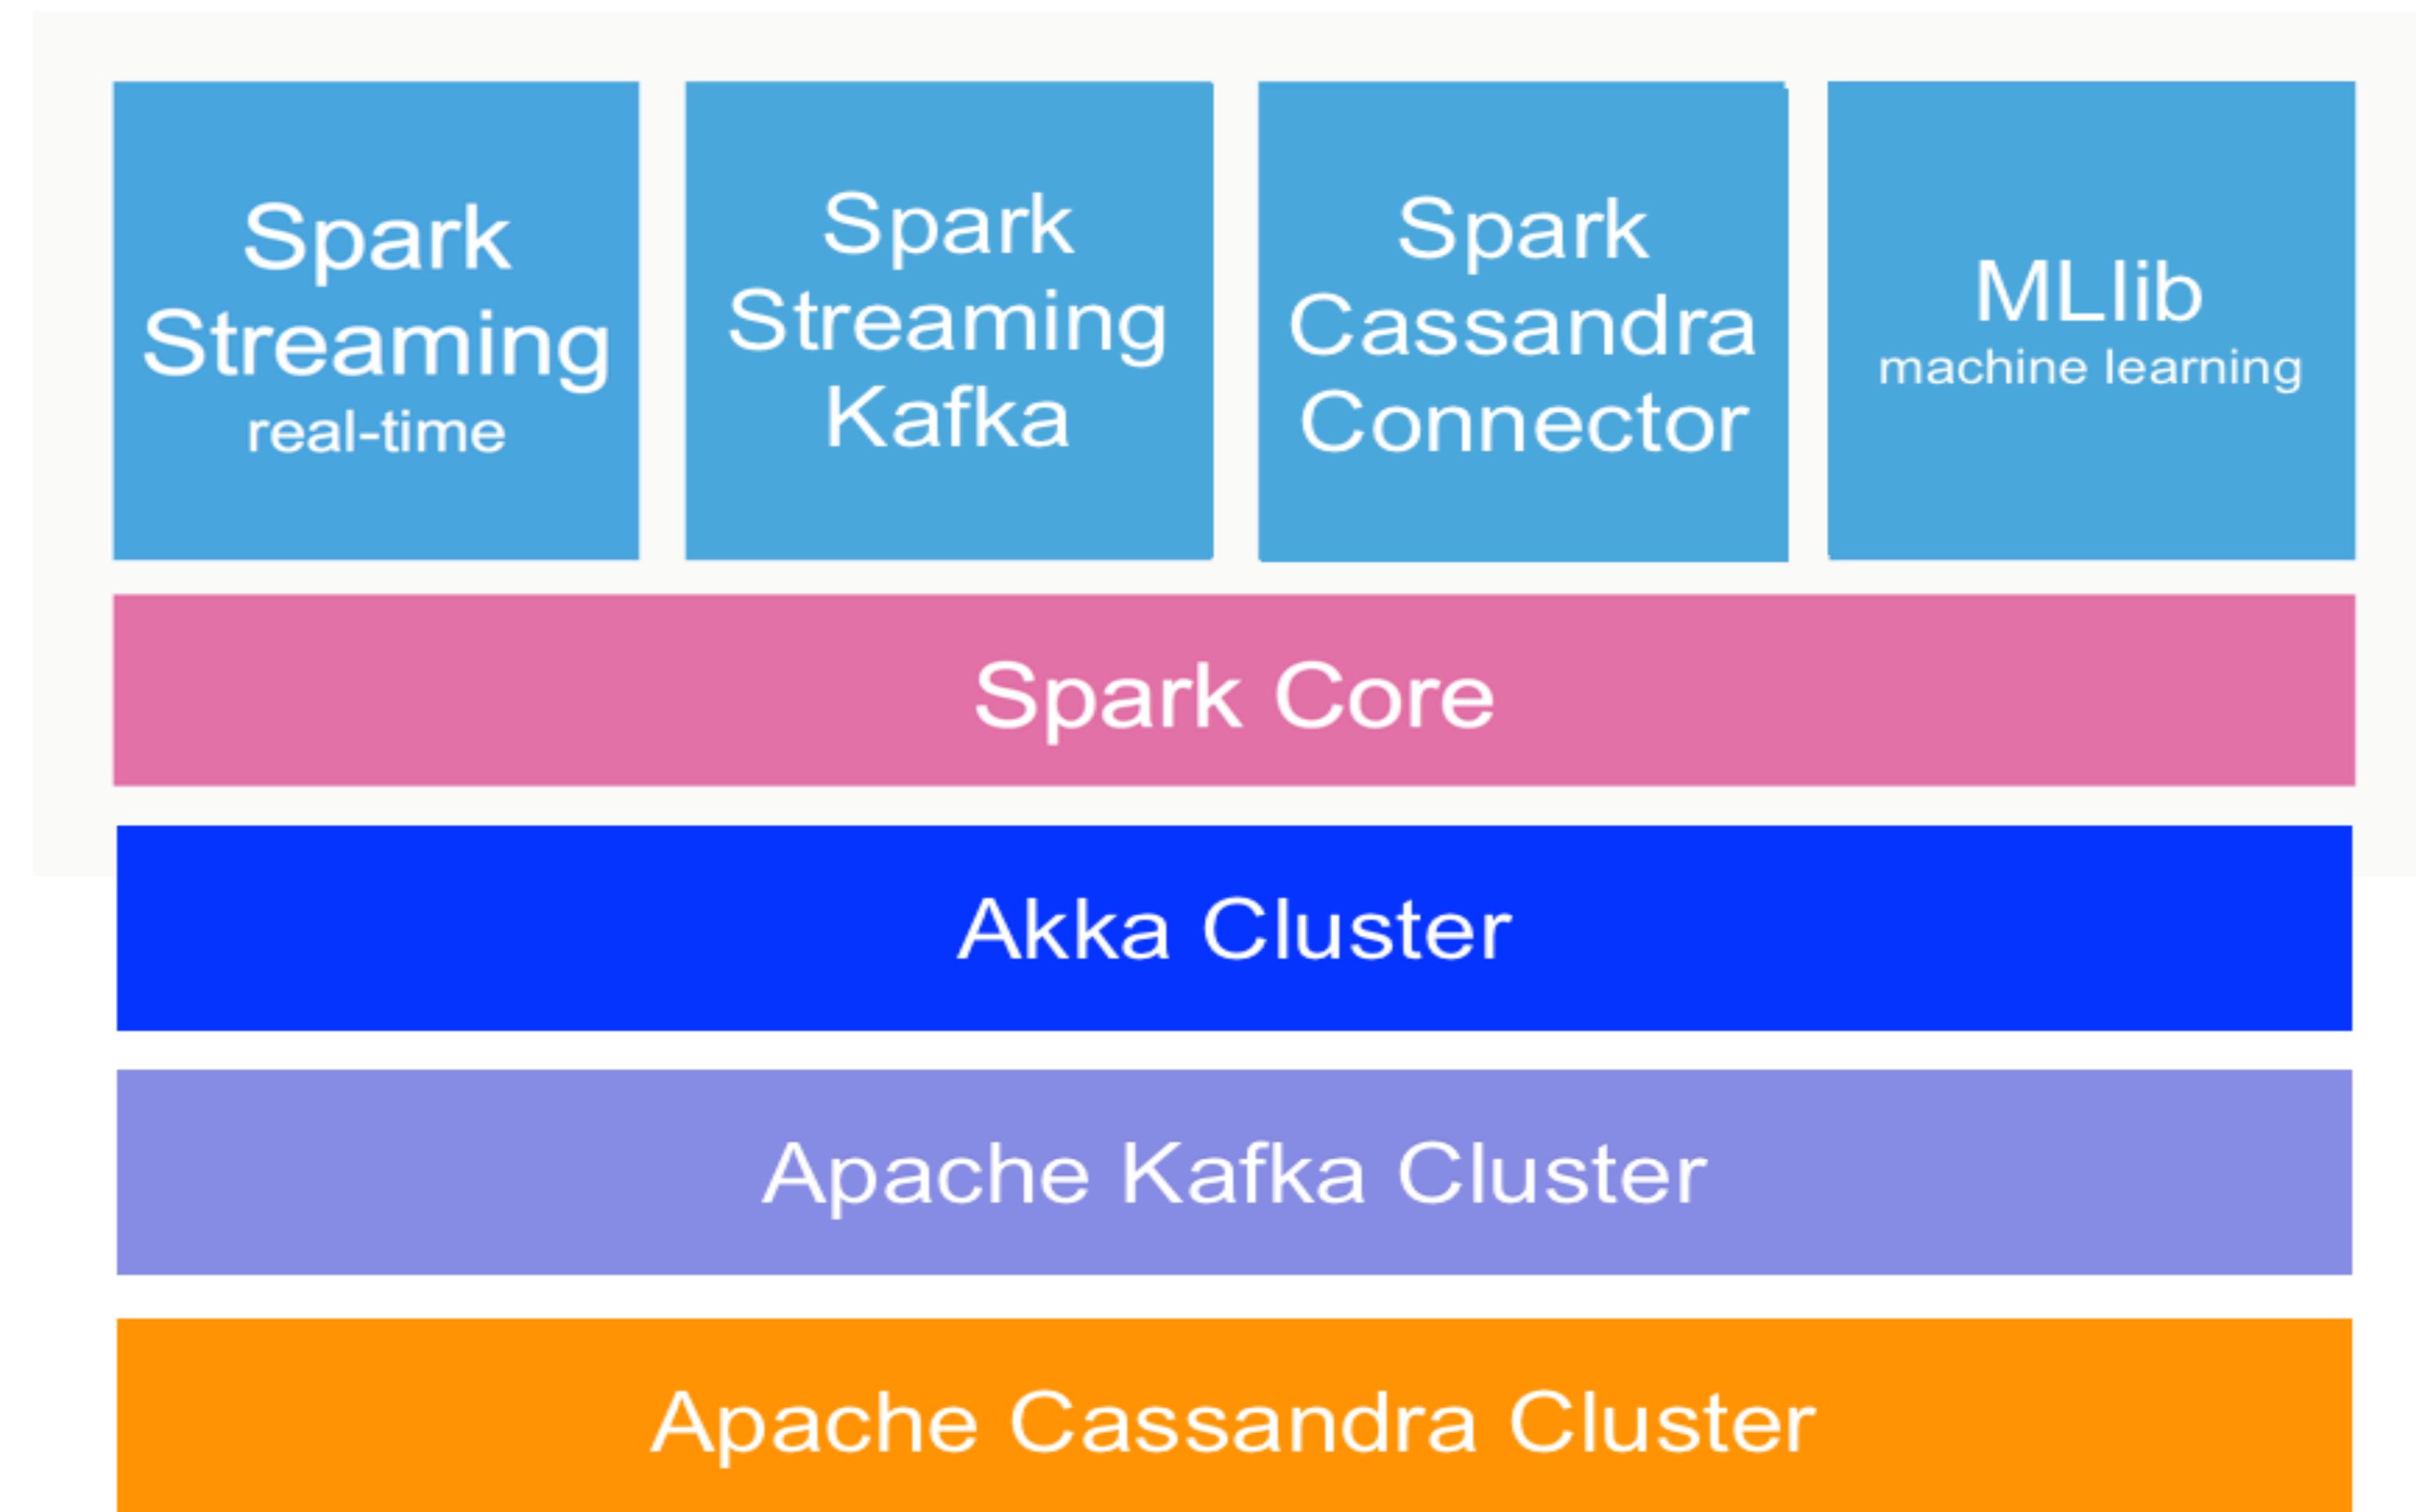
What We Need

- Fault Tolerant
- Failure Detection
- Fast - low latency, distributed, data locality
- Masterless, Decentralized Cluster Membership
- Span Racks and DataCenters
 - Hashes The Node Ring
 - Partition-Aware
- Elasticity
- Asynchronous - message-passing system
- Parallelism
- Network Topology Aware

Lambda Architecture with Spark, Kafka, Cassandra and Akka (Scala!)

Lambda Architecture - is a data-processing architecture designed to handle massive quantities of data by taking advantage of both batch and stream processing methods.

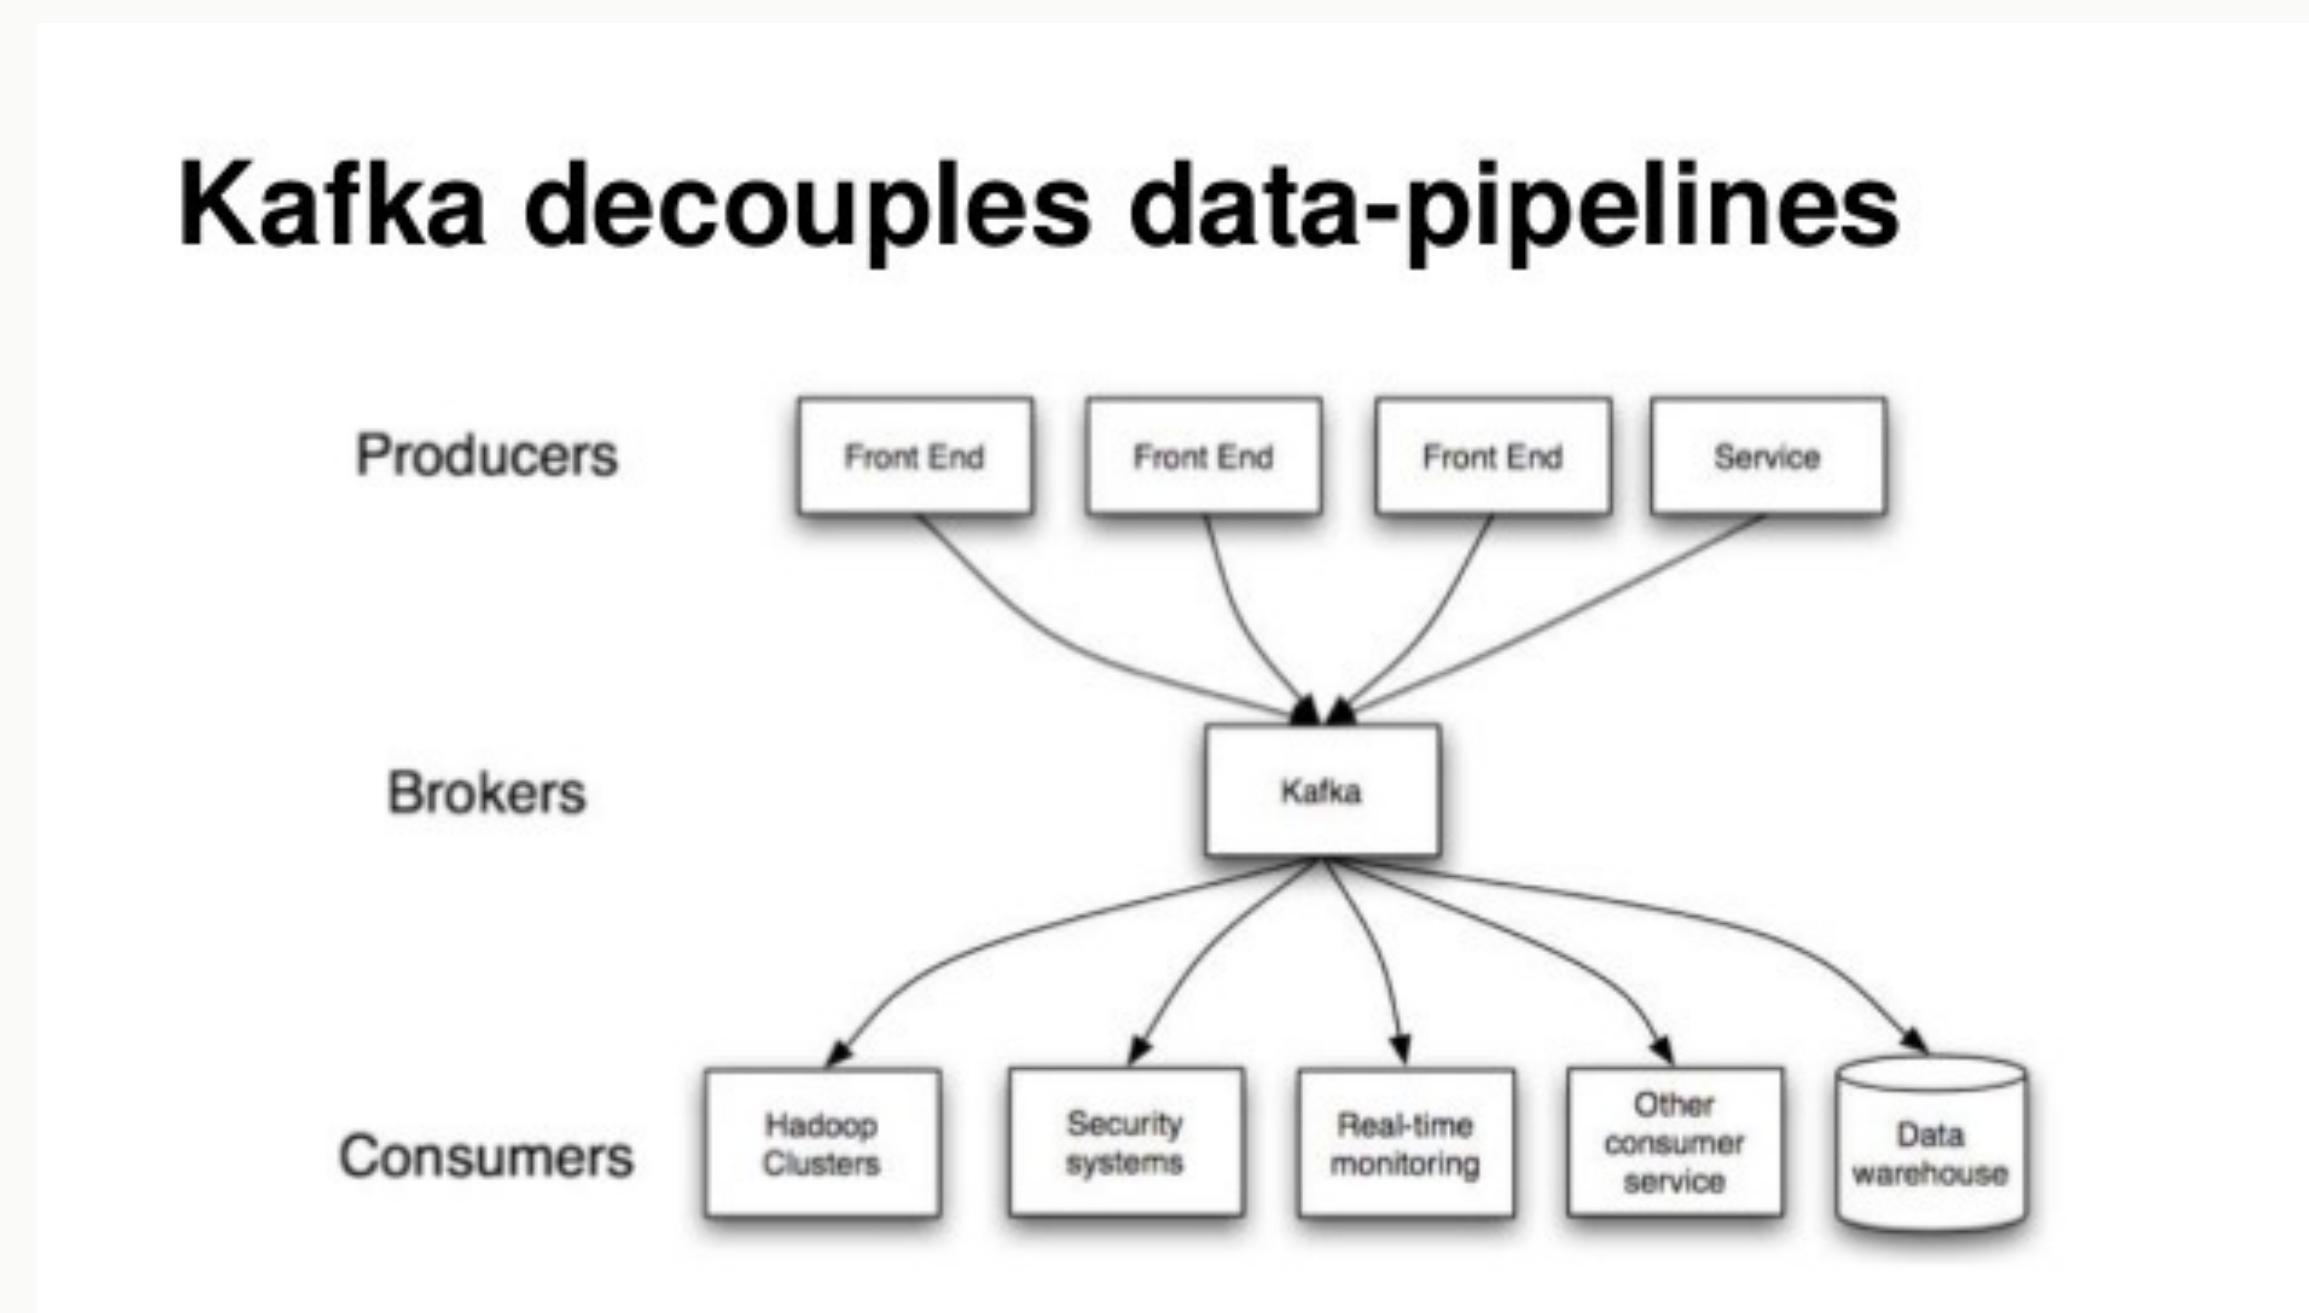
- Spark - one of the few, if not the only, data processing framework that allows you to have both batch and stream processing of terabytes of data in the same application.
- Storm does not



Why Akka Rocks

- location transparency
- fault tolerant
- async message passing
- non-deterministic
- share nothing
- actor atomicity (w/in actor)

Apache Kafka



From Joe Stein

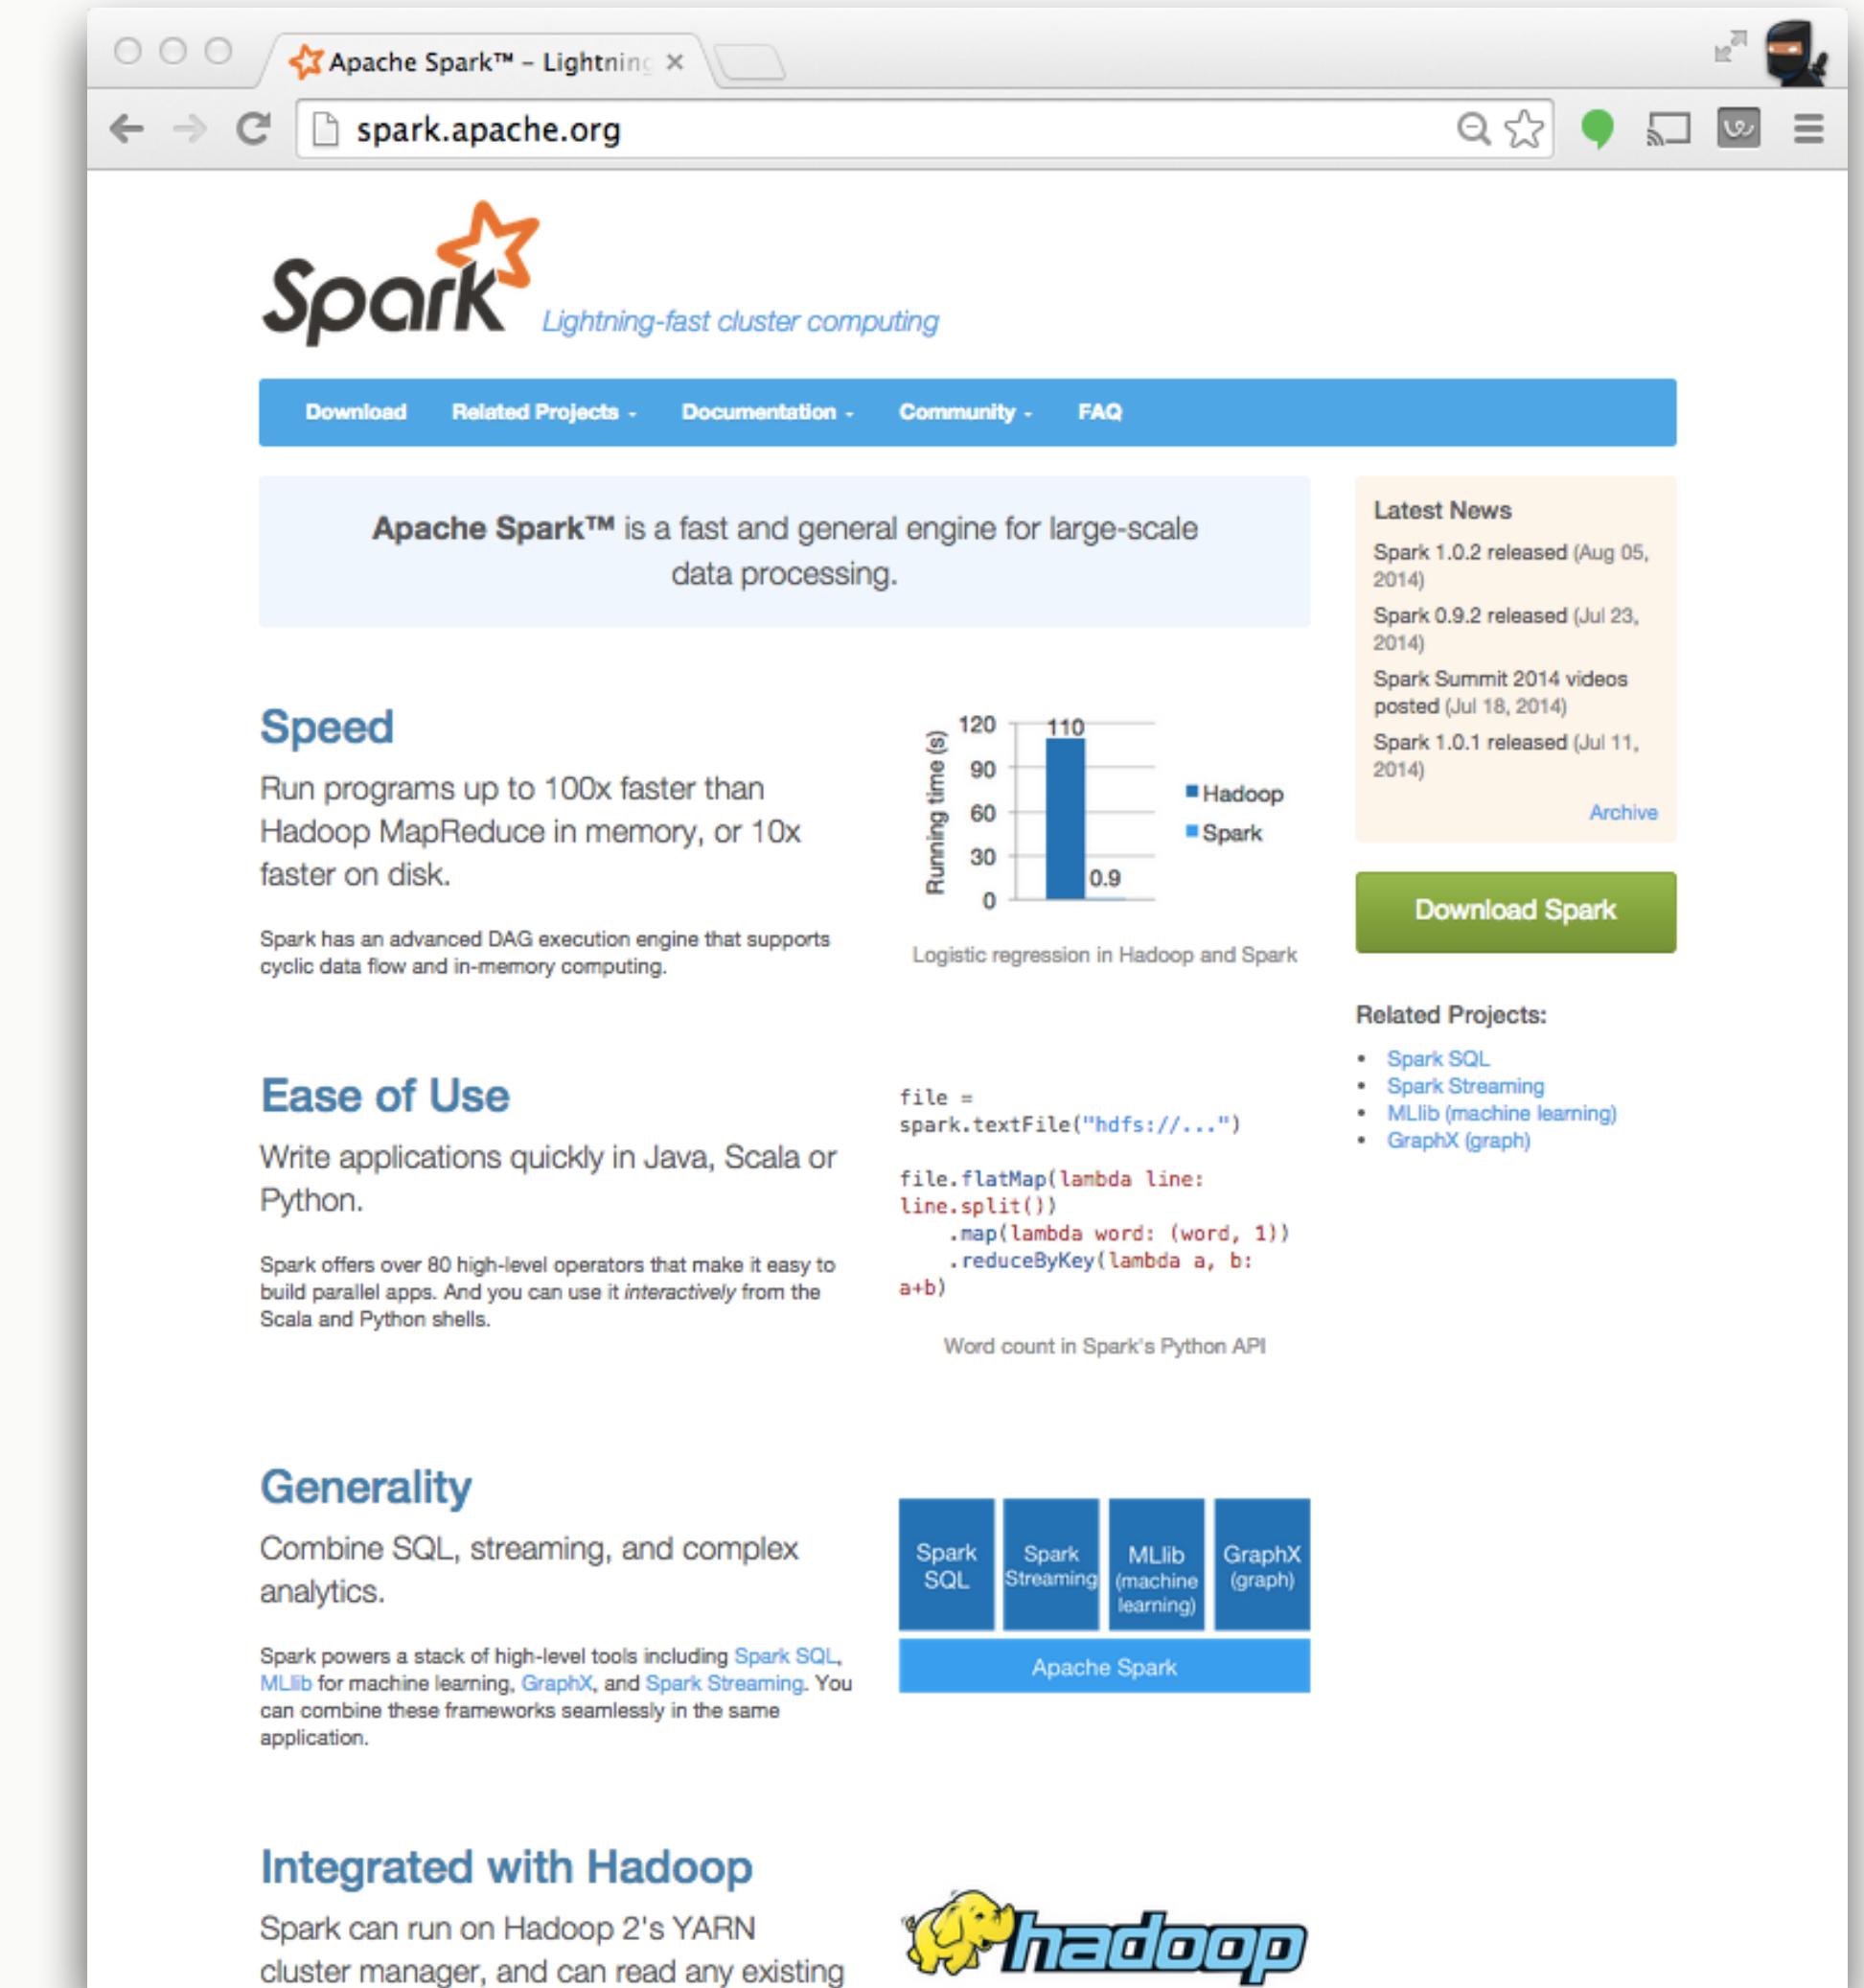
Apache Spark and Spark Streaming

The Drive-By Version

DATASTAX

What Is Apache Spark

- Fast, distributed, scalable and fault tolerant cluster compute system
- Enables Low-latency with complex analytics
- Developed in 2009 at UC Berkeley AMPLab, open sourced in 2010, and became a top-level Apache project in February, 2014



The screenshot shows the official Apache Spark website at spark.apache.org. The page features a prominent logo with the word "Spark" and a red starburst icon, followed by the tagline "Lightning-fast cluster computing". A navigation bar includes links for Download, Related Projects, Documentation, Community, and FAQ. On the left, a main content area highlights "Apache Spark™ is a fast and general engine for large-scale data processing". It features sections on Speed (comparing Hadoop and Spark execution times), Ease of Use (showing Scala code for word count), Generality (mentioning integration with SQL, streaming, and machine learning), and Integrated with Hadoop (mentioning YARN support). To the right, there's a "Latest News" sidebar with links to recent releases and a "Related Projects" sidebar listing Spark SQL, Spark Streaming, MLlib, and GraphX. A "Download Spark" button is also visible.

Speed

Run programs up to 100x faster than Hadoop MapReduce in memory, or 10x faster on disk.

Spark has an advanced DAG execution engine that supports cyclic data flow and in-memory computing.

	Hadoop	Spark
Running time (s)	110	0.9

Logistic regression in Hadoop and Spark

Ease of Use

```
file = spark.textFile("hdfs://...")  
file.flatMap(lambda line:  
    line.split()  
    .map(lambda word:(word, 1))  
    .reduceByKey(lambda a, b:  
        a+b))
```

Word count in Spark's Python API

Generality

Combine SQL, streaming, and complex analytics.

Spark powers a stack of high-level tools including [Spark SQL](#), [MLlib](#) for machine learning, [GraphX](#), and [Spark Streaming](#). You can combine these frameworks seamlessly in the same application.

Integrated with Hadoop

Spark can run on Hadoop 2's YARN cluster manager, and can read any existing

 hadoop

Most Active OSS In Big Data



apache / spark
mirrored from <git://git.apache.org/spark.git>

Watch ▾ 427 Star 2,285 Fork 2,044

November 11, 2014 – December 11, 2014 Period: 1 month ▾

Overview

154 Active Pull Requests	0 Active Issues
10 Merged Pull Requests	0 Closed Issues
154 Proposed Pull Requests	0 New Issues

Excluding merges, **317 authors** have pushed **297 commits** to master and **3,546 commits** to all branches. On master, **702 files** have changed and there have been **38,458 additions** and **18,021 deletions**.

A bar chart showing the number of pull requests proposed by 154 different people. The y-axis represents the count of pull requests, ranging from 0 to 150. The x-axis lists 154 individuals, each represented by a small profile picture. The bars are orange.

154 Pull requests proposed by 82 people

Apache Spark - Easy to Use & Fast

- 10x faster on disk, 100x faster in memory than Hadoop MR
- Write, test and maintain 2 - 5x less code
- Fault Tolerant Distributed Datasets
- Batch, iterative and streaming analysis
- In Memory Storage and Disk
- Integrates with Most File and Storage Options

Easy API

- Abstracts complex algorithms to high level functions
- Collections API over large datasets
- Uses a functional programming model - clean
- Scala, Java and Python APIs
- Use from a REPL
- Easy integration with SQL, streaming, ML, Graphing, R...

Resilient Distributed Dataset (RDD)

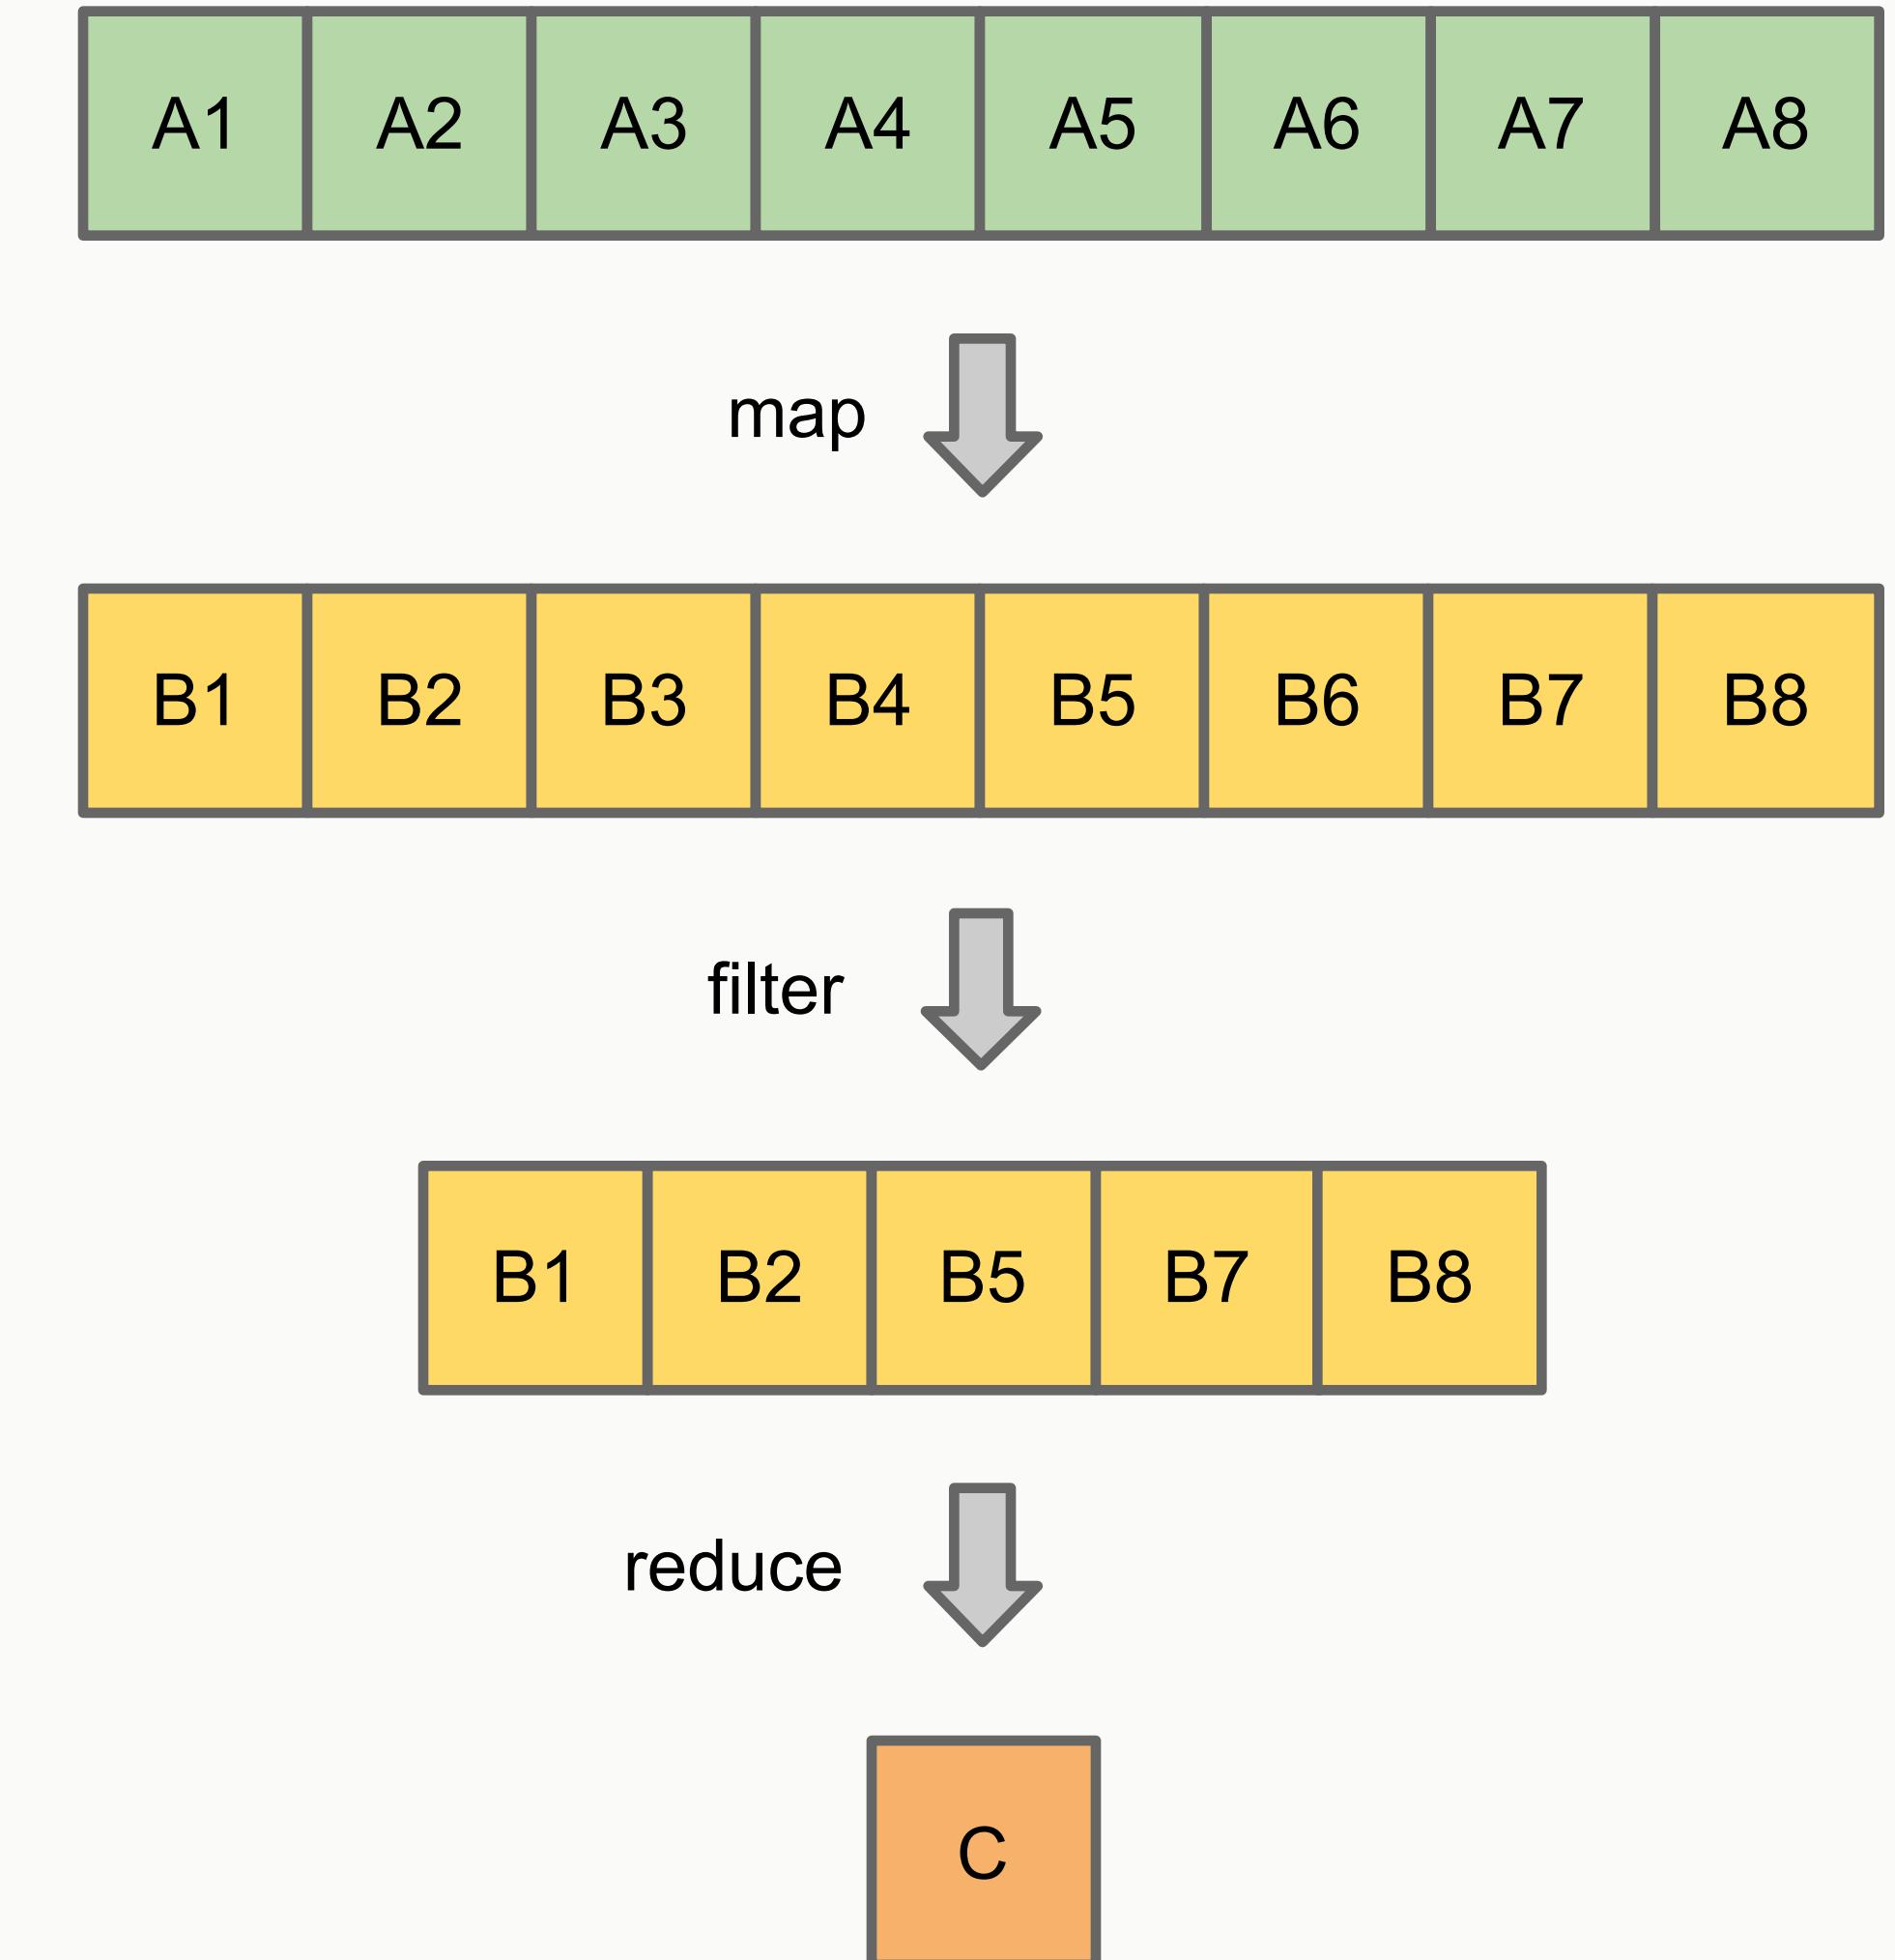
- Fault Tolerant: can recalculate from any point of failure
- Created through transformations on data (map, filter..) or other RDDs
- Immutable
- Partitioned
 - Indicate which RDD elements to partition across machines based on a key in each record
- Can be reused

Spark Data Model

Resilient Distributed Dataset

A collection:

- immutable
- iterable
- serializable
- distributed
- parallel
- lazy



RDD Operations

```
16  object SparkWordCount extends WordCountBlueprint {  
17  
18      sc.textFile("./src/main/resources/data/words")  
19          .flatMap(_.split("\\s+"))  
20          .map(word => (clean(word), 1))  
21          .reduceByKey(_ + _)  
22          .collect foreach println |  
23  }
```

Transformation

Action

Collection To RDD

```
scala> val data = Array(1, 2, 3, 4, 5)
data: Array[Int] = Array(1, 2, 3, 4, 5)
```

```
scala> val distData = sc.parallelize(data)
distData: spark.RDD[Int] =
spark.ParallelCollection@10d13e3e
```

Spark Basic Word Count

```
val conf = new SparkConf()
    .setMaster("local[*]")
    .setAppName("Simple Word Count")

val sc = new SparkContext(conf)

sc.textFile(words)
    .flatMap(_.split("\\s+"))
    .map(word => (word.toLowerCase, 1))
    .reduceByKey(_ + _)
    .collect foreach println
```

Apache Spark - Easy to Use API

```
/** Returns the top (k) highest temps for any location in the `year`. */
val k = 20
```

```
def topK(aggregate: Seq[Double]): Seq[Double] =
  sc.parallelize(aggregate).top(k)
```

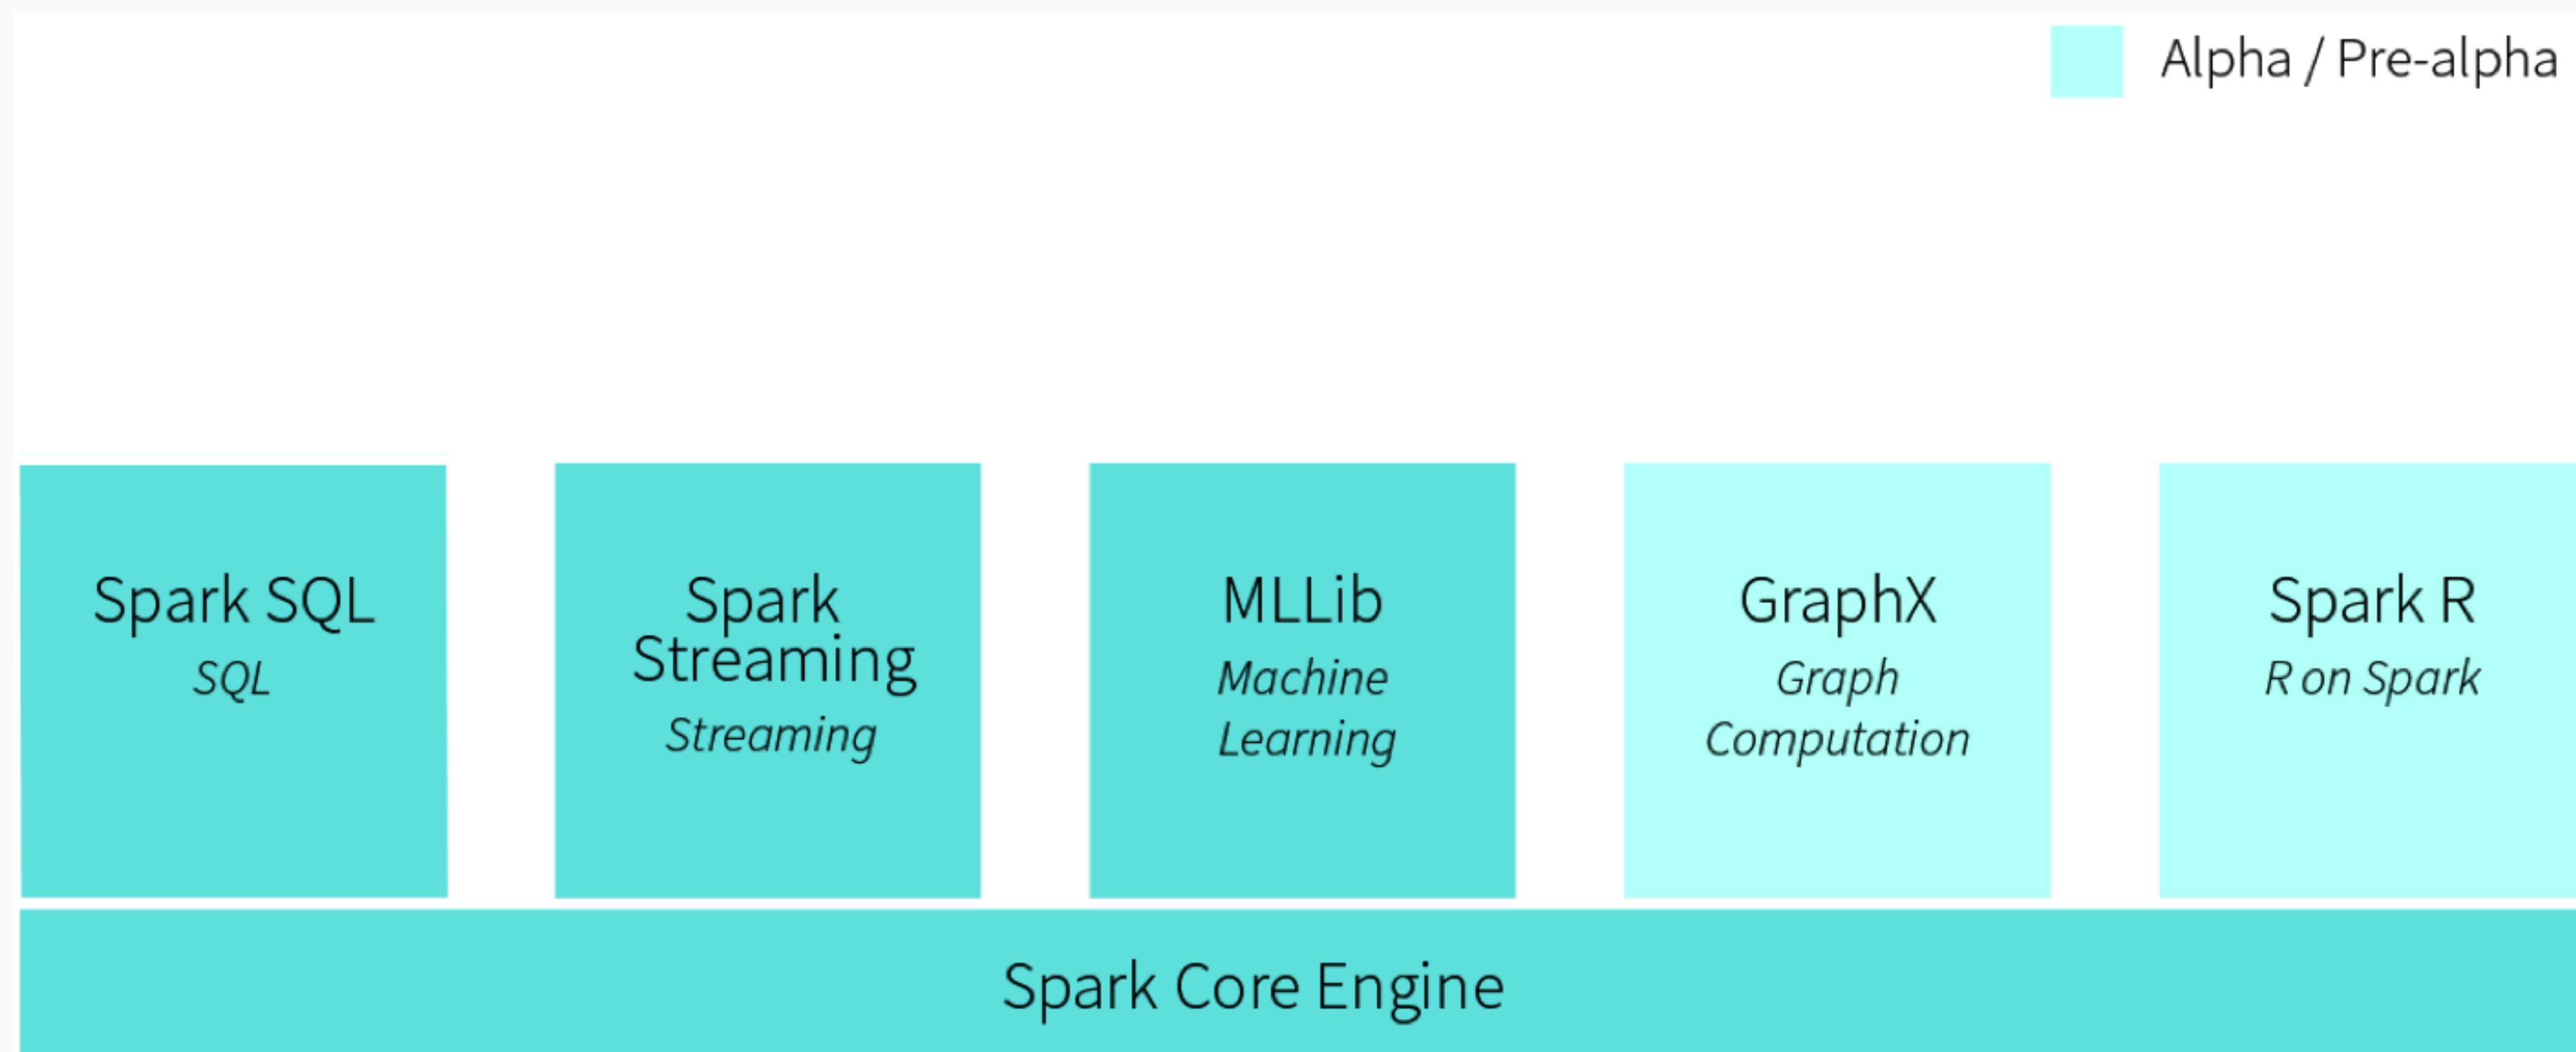
```
/** Returns the top (k) highest temps ... in a Future */
def topK(aggregate: Seq[Double]): Future[Seq[Double]] =
  sc.parallelize(aggregate).top(k).collectAsync
```

Not Just MapReduce



<pre>def aggregate[U](zeroValue: U)(seqOp: (U, T) → U, combOp: (U, U) → U)(implicit arg0: ClassTag[U]): RDD[U]</pre> <p>Aggregate the elements of each partition, and then the results for all the partitions.</p>	<pre>def glom(): RDD[Array[T]]</pre> <p>Return an RDD created by coalescing all elements within each partition into an array.</p>	<pre>def min()(implicit ord: Ordering[T]): T</pre> <p>Returns the min of this RDD as defined by the implicit Ordering[T].</p>
<pre>def cache(): RDD.this.type</pre> <p>Persist this RDD with the default storage level (MEMORY_ONLY).</p>	<pre>def groupBy[K](f: (T) → K, p: Partitioner)(implicit kt: KeyType[K]): RDD[K]</pre> <p>Return an RDD of grouped items.</p>	<pre>var name: String</pre> <p>A friendly name for this RDD.</p>
<pre>def cartesian[U](other: RDD[U])(implicit arg0: ClassTag[U]): RDD[(U, U)]</pre> <p>Return the Cartesian product of this RDD and another one, that is, the product of all elements of this RDD with all elements of another one.</p>	<pre>def groupBy[K](f: (T) → K, numPartitions: Int)(implicit kt: KeyType[K]): RDD[K]</pre> <p>Return an RDD of grouped elements.</p>	<pre>val partitioner: Option[Partitioner]</pre> <p>Optionally overridden by subclasses to specify how they are partitioned.</p>
<pre>def checkpoint(): Unit</pre> <p>Mark this RDD for checkpointing.</p>	<pre>def groupBy[K](f: (T) → K)(implicit kt: ClassTag[K]): RDD[K]</pre> <p>Return an RDD of grouped items.</p>	<pre>def partitions: Array[Partition]</pre> <p>Get the array of partitions of this RDD, taking into account whether the RDD is checked for consistency.</p>
<pre>def coalesce(numPartitions: Int, shuffle: Boolean = false): RDD[T]</pre> <p>Return a new RDD that is reduced into numPartitions partitions.</p>	<pre>val id: Int</pre> <p>A unique ID for this RDD (within its SparkContext).</p>	<pre>def persist(): RDD.this.type</pre> <p>Persist this RDD with the default storage level (MEMORY_ONLY).</p>
<pre>def collect[U](f: PartialFunction[T, U])(implicit arg0: ClassTag[U]): RDD[U]</pre> <p>Return an RDD that contains all matching values by applying f.</p>	<pre>def intersection(other: RDD[T], numPartitions: Int): RDD[T]</pre> <p>Return the intersection of this RDD and another one.</p>	<pre>def persist(newLevel: StorageLevel): RDD.this.type</pre> <p>Set this RDD's storage level to persist its values across operations after the first time it is computed.</p>
<pre>def collect(): Array[T]</pre> <p>Return an array that contains all of the elements in this RDD.</p>	<pre>def intersection(other: RDD[T], partitioner: Partitioner): RDD[T]</pre> <p>Return the intersection of this RDD and another one.</p>	<pre>def pipe(command: Seq[String], env: Map[String, String] = Map()): RDD[String]</pre> <p>Return an RDD created by piping elements to a forked external process.</p>
<pre>def context: SparkContext</pre> <p>The org.apache.spark.SparkContext that this RDD was created on.</p>	<pre>def intersection(other: RDD[T]): RDD[T]</pre> <p>Return the intersection of this RDD and another one.</p>	<pre>def pipe(command: String, env: Map[String, String]): RDD[String]</pre> <p>Return an RDD created by piping elements to a forked external process.</p>
<pre>def count(): Long</pre> <p>Return the number of elements in the RDD.</p>	<pre>def isCheckpointed: Boolean</pre> <p>Return whether this RDD has been checkpointed or not.</p>	<pre>def pipe(command: String): RDD[String]</pre> <p>Return an RDD created by piping elements to a forked external process.</p>
<pre>def countApprox(timeout: Long, confidence: Double = 0.95): Long</pre> <p>Approximate version of count() that returns a potentially incomplete result.</p>	<pre>def iterator(split: Partition, context: TaskContext): Iterator[T]</pre> <p>Internal method to this RDD; will read from cache if applicable, or otherwise read from disk.</p>	<pre>def preferredLocations(split: Partition): Seq[String]</pre> <p>Get the preferred locations of a partition (as hostnames), taking into account whether the partition is replicated.</p>
<pre>def countApproxDistinct(relativeSD: Double = 0.05): Long</pre> <p>Return approximate number of distinct elements in the RDD.</p>	<pre>def keyBy[K](f: (T) → K): RDD[(K, T)]</pre> <p>Creates tuples of the elements in this RDD by applying f.</p>	<pre>def randomSplit(weights: Array[Double], seed: Long = Utils.random.nextInt()): RDD[Dataset[_]]</pre> <p>Randomly splits this RDD with the provided weights.</p>
<pre>def countApproxDistinct(p: Int, sp: Int): Long</pre> <p>Return approximate number of distinct elements in the RDD.</p>	<pre>def map[U](f: (T) → U)(implicit arg0: ClassTag[U]): RDD[U]</pre> <p>Return a new RDD by applying a function to all elements of this RDD.</p>	<pre>def reduce(f: (T, T) → T): T</pre> <p>Reduces the elements of this RDD using the specified commutative and associative reduce function.</p>
<pre>def countByValue()(implicit ord: Ordering[T] = null): Map[T, Long]</pre> <p>Return the count of each unique value in this RDD as a map of (value, count).</p>	<pre>def mapPartitions[U](f: (Iterator[T]) → Iterator[U], p: Partitioner): RDD[U]</pre> <p>Return a new RDD by applying a function to each partition of this RDD.</p>	<pre>def repartition(numPartitions: Int)(implicit ord: Ordering[T] = null): RDD[T]</pre> <p>Return a new RDD that has exactly numPartitions partitions.</p>
<pre>def countByValueApprox(timeout: Long, confidence: Double): Map[T, Long]</pre> <p>Approximate version of countByValue().</p>	<pre>def mapPartitionsWithContext[U](f: (TaskContext, Iterator[T]) → Iterator[U], p: Partitioner): RDD[U]</pre> <p>Return a new RDD by applying a function to each partition of this RDD.</p>	<pre>def sample(withReplacement: Boolean, fraction: Double, seed: Long): RDD[T]</pre> <p>Return a sampled subset of this RDD.</p>
<pre>def dependencies: Seq[Dependency[_]]</pre> <p>Get the list of dependencies of this RDD, taking into account whether the RDD is checked for consistency.</p>	<pre>def mapPartitionsWithIndex[U](f: (Int, Iterator[T]) → Iterator[U], p: Partitioner): RDD[U]</pre> <p>Return a new RDD by applying a function to each partition of this RDD.</p>	<pre>def saveAsObjectFile(path: String): Unit</pre> <p>Save this RDD as a SequenceFile of serialized objects.</p>
<pre>def distinct(): RDD[T]</pre> <p>Return a new RDD containing the distinct elements in this RDD.</p>	<pre>def max()(implicit ord: Ordering[T]): T</pre> <p>Returns the max of this RDD as defined by the implicit Ordering[T].</p>	<pre>def saveAsTextFile(path: String, codec: Class[_ <: CompressionCodec]): Unit</pre> <p>Save this RDD as a compressed text file, using string representations of elements.</p>

Spark Components



Why Scala?

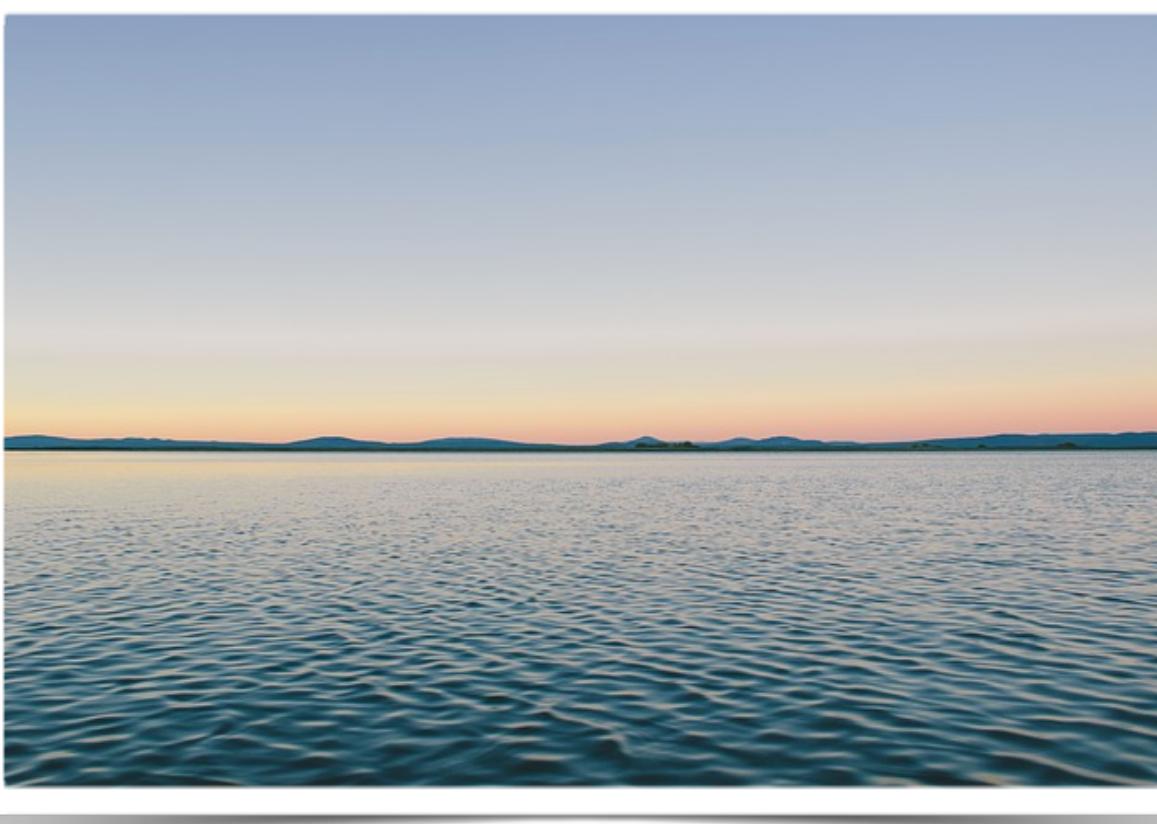
<http://apache-spark-user-list.1001560.n3.nabble.com/Why-Scala-tp6536p6538.html>

- Functional
- On the JVM
- Capture functions and ship them across the network
- Static typing - easier to control performance
- Leverage REPL Spark REPL

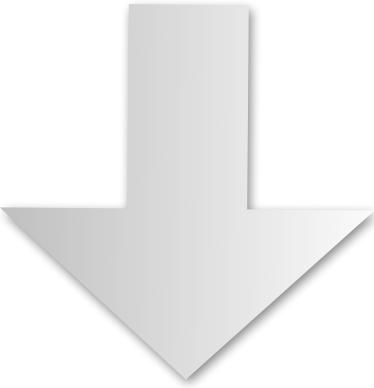
When Batch Is Not Enough

- Scheduled batch only gets you so far
- I want results continuously in the event stream
- I want to run computations in my even-driven asynchronous apps

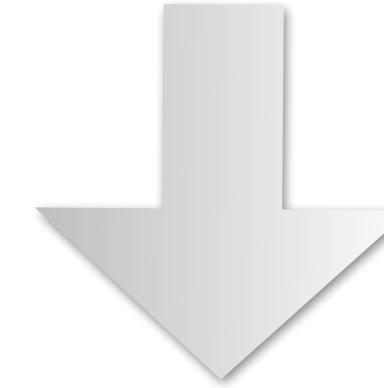
Spark Streaming



zillions of bytes



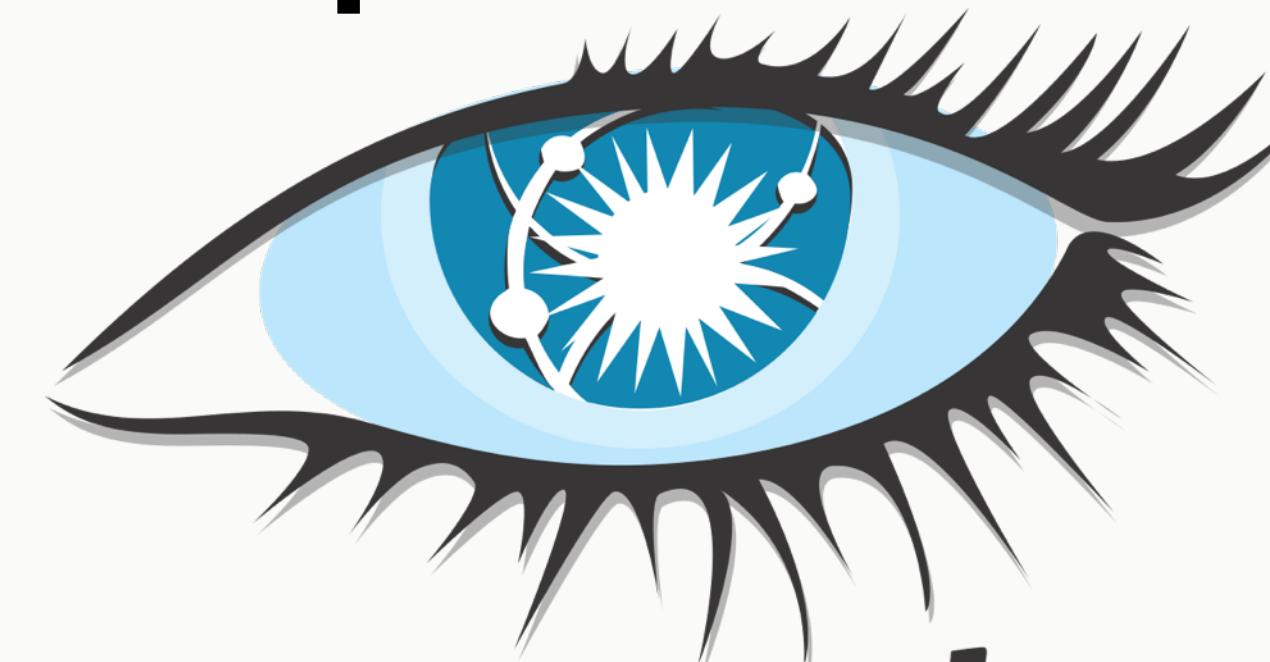
gigabytes per second



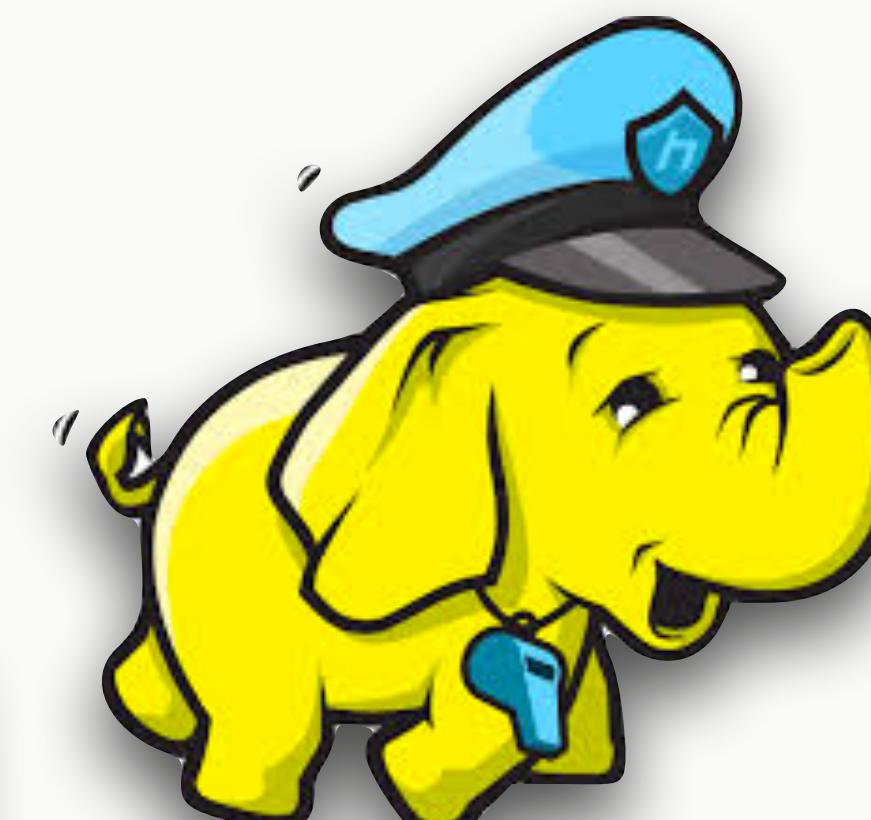
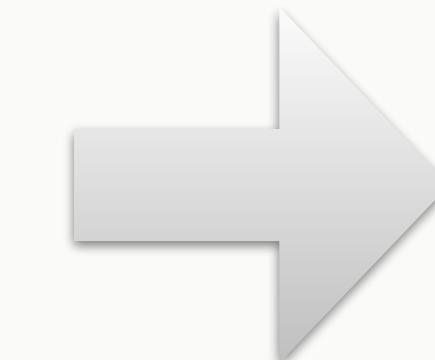
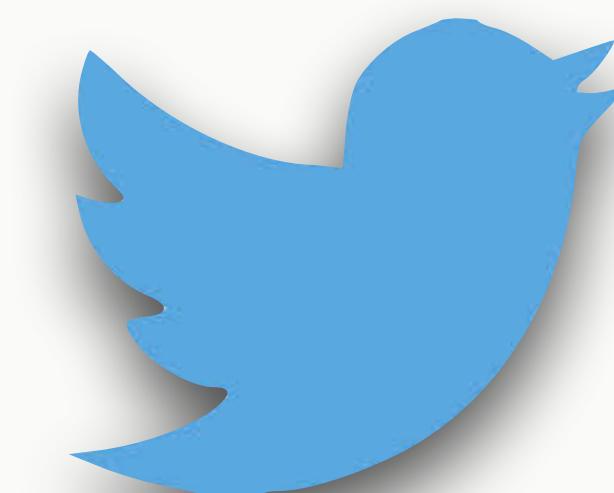
Input & Output Sources



akka

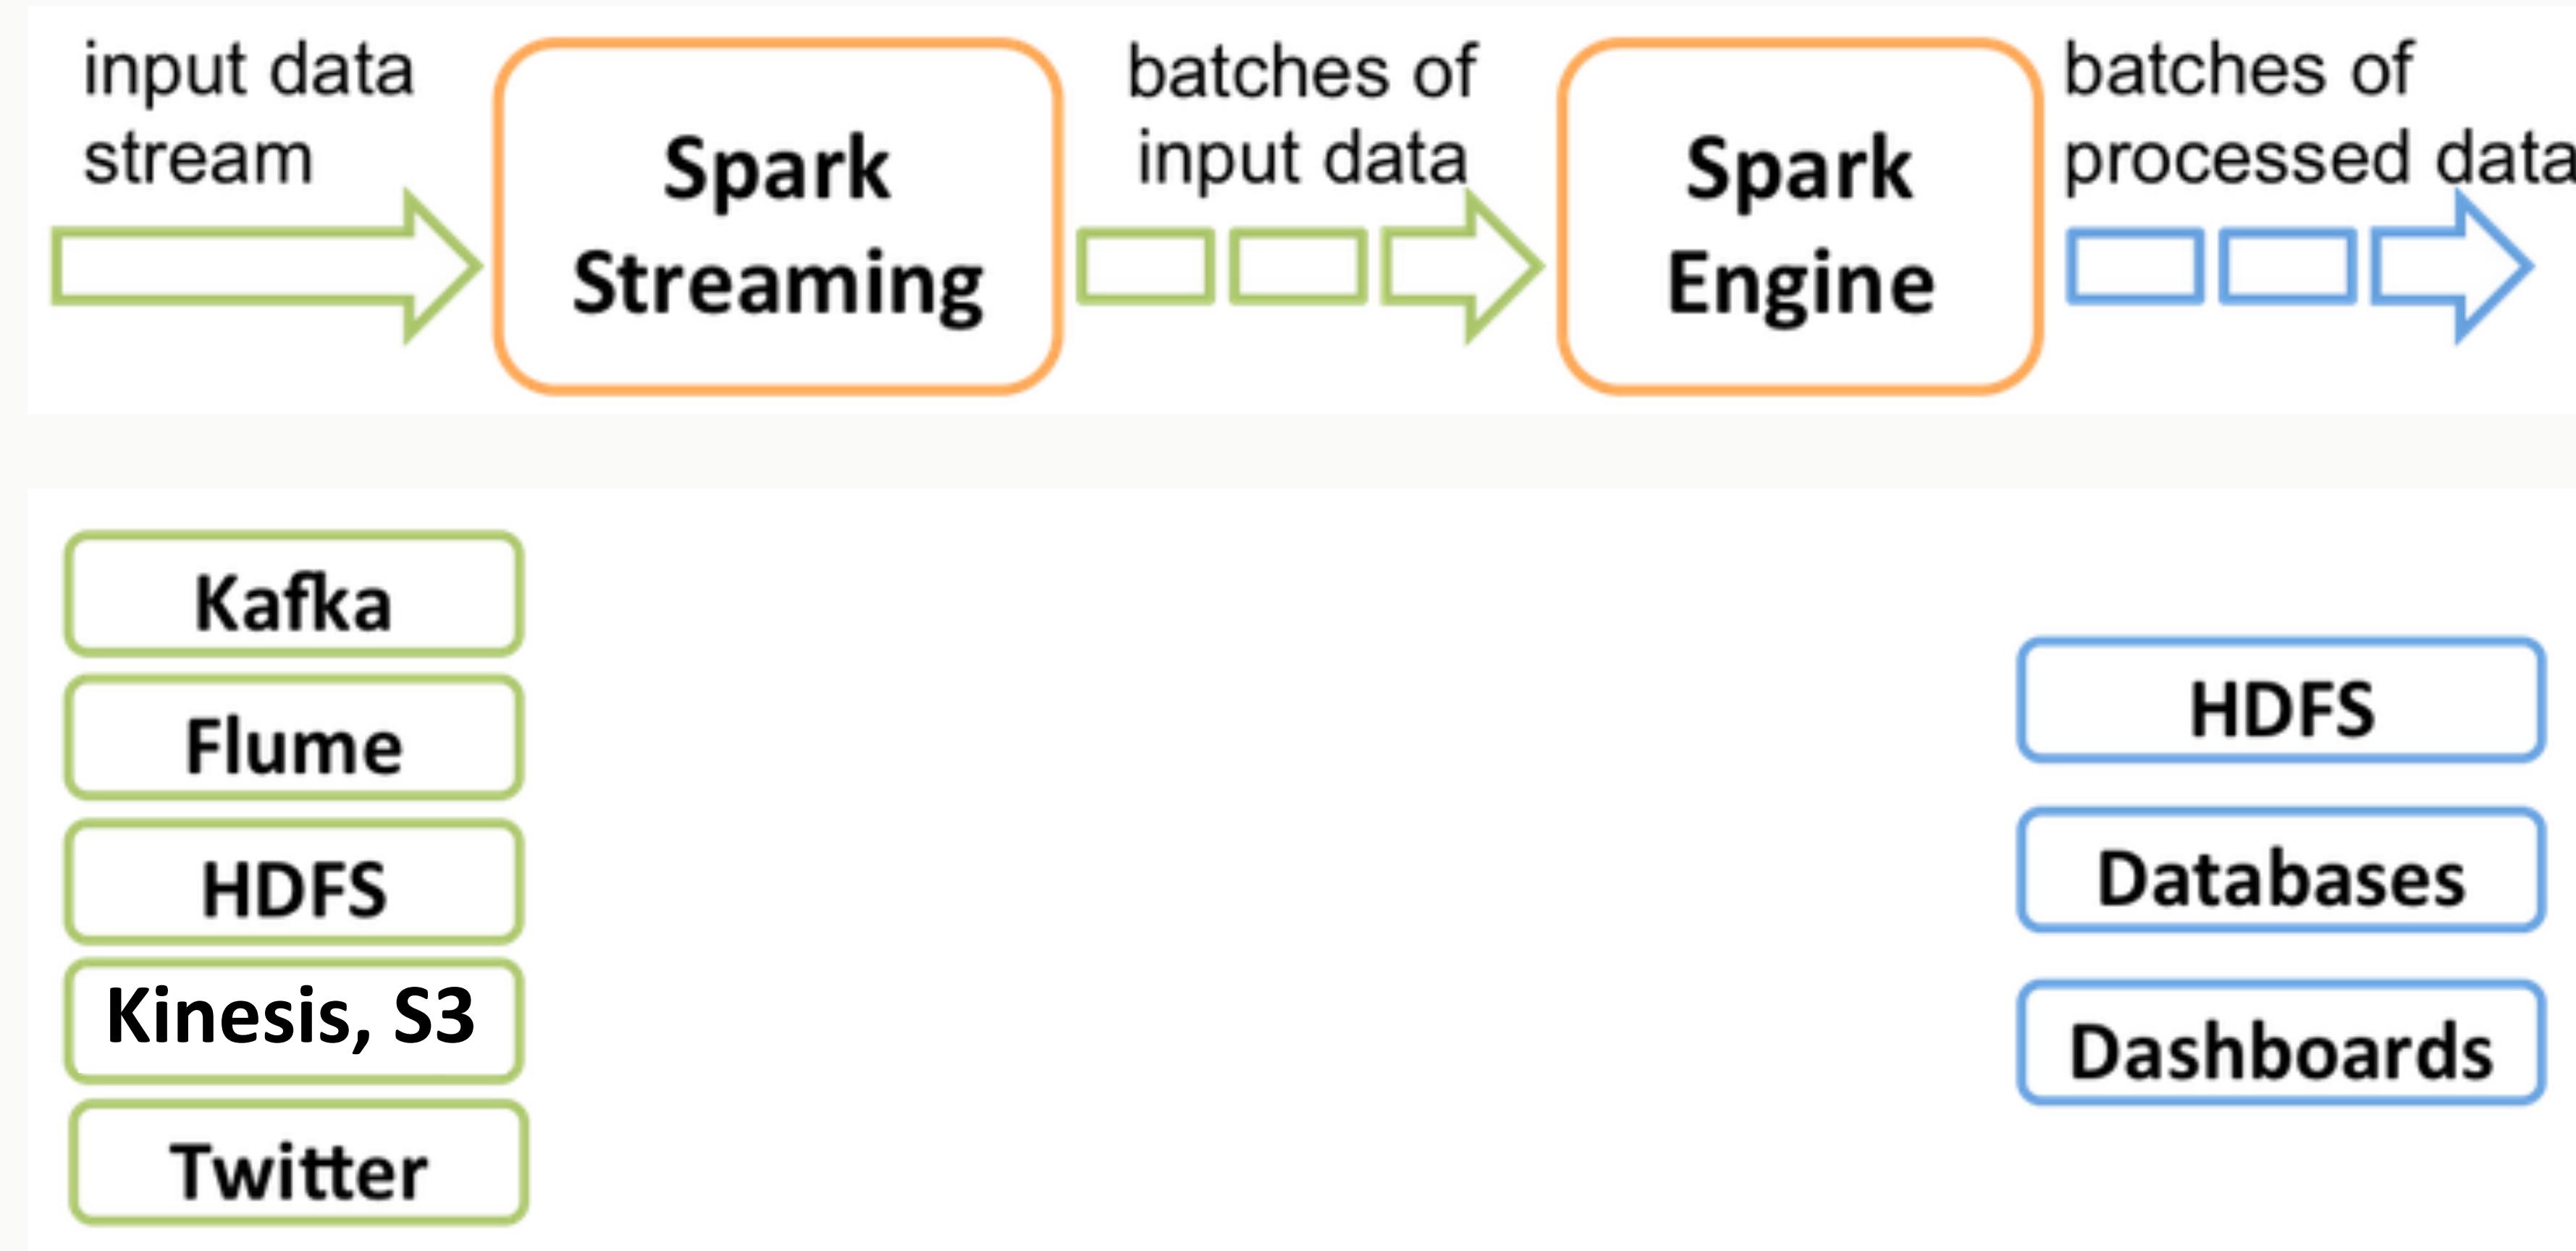


cassandra



ØMQ

Spark Streaming



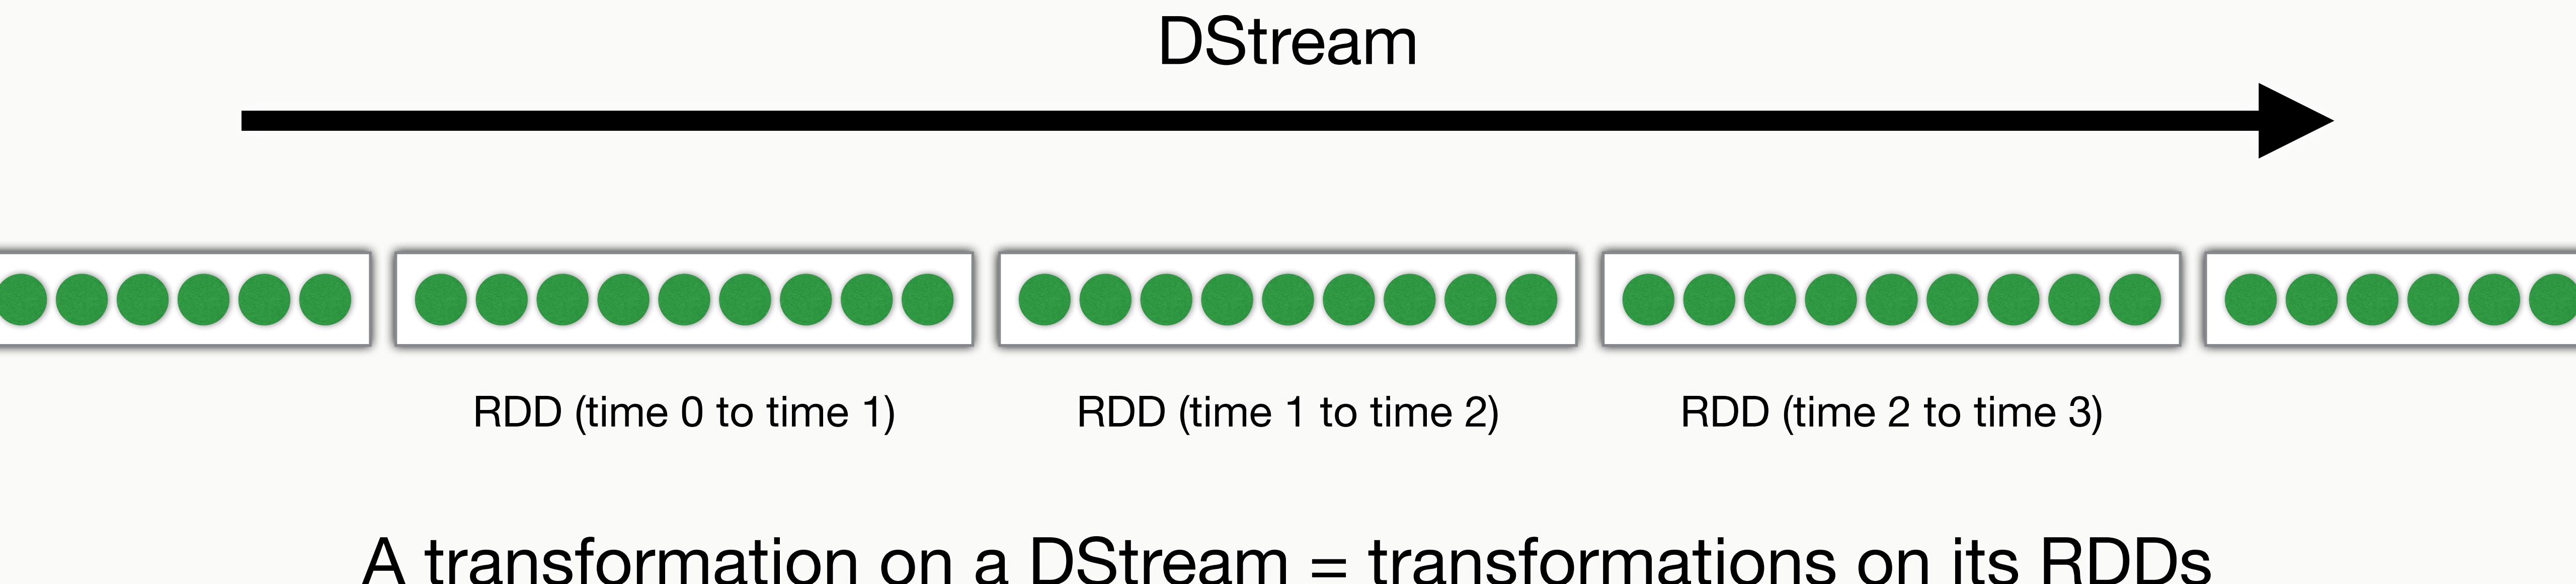
Common Use Cases

applications	sensors	web	mobile phones
intrusion detection	malfunction detection	site analytics	network metrics analysis
fraud detection	dynamic process optimisation	recommendations	location based ads
log processing	supply chain planning	sentiment analysis	...

DStream (Discretized Stream)

Continuous stream of micro batches

- Complex processing models with minimal effort
- Streaming computations on small time intervals



A transformation on a DStream = transformations on its RDDs

InputDStreams and Receivers

DStreams - the stream of raw data received from streaming sources:

- Basic Source - in the StreamingContext API
- Advanced Source - in external modules and separate Spark artifacts

Receivers

- Reliable Receivers - for data sources supporting acks (like Kafka)
- Unreliable Receivers - for data sources not supporting acks

Spark Streaming Modules

GroupId	ArtifactId	Latest Version
org.apache.spark	spark-streaming-kinesis-asl_2.10	1.1.0
org.apache.spark	spark-streaming-mqtt_2.10	1.1.0 all (7)
org.apache.spark	spark-streaming-zeromq_2.10	1.1.0 all (7)
org.apache.spark	spark-streaming-flume_2.10	1.1.0 all (7)
org.apache.spark	spark-streaming-flume-sink_2.10	1.1.0
org.apache.spark	spark-streaming-kafka_2.10	1.1.0 all (7)
org.apache.spark	spark-streaming-twitter_2.10	1.1.0 all (7)

Spark Streaming Setup

```
val conf = new SparkConf().setMaster(SparkMaster).setAppName(AppName)
.set("spark.serializer", "org.apache.spark.serializer.KryoSerializer")
.set("spark.kryo.registrator", "com.datastax.killrweather.KillrKryoRegistrator")
.set("spark.cleaner.ttl", SparkCleanerTtl)

val streamingContext = new StreamingContext(conf, Milliseconds(500))

// Do work in the stream

ssc.checkpoint(checkpointDir)
ssc.start()
ssc.awaitTermination
```

Checkpointing

Allows saving enough of information to a fault-tolerant storage to allow the RDDs

- Metadata - the information defining the streaming computation
- Data (RDDs)

Usage

- With `updateStateByKey`, `reduceByKeyAndWindow` - stateful transformations
- To recover from failures in Spark Streaming apps

Can affect performance, depending on

- The data and or batch sizes.
- The speed of the file system that is being used for checkpointing

Basic Streaming: FileInputStream



```
// Creates new DStreams
ssc.textFileStream("s3n://raw_data_bucket/")
  .flatMap(_.split("\\s+"))
  .map(_.toLowerCase, 1))
  .countByValue()
  .saveAsObjectFile("s3n://analytics_bucket/")
```

ReceiverInputDStreams

```
val stream = KafkaUtils.createStream[String, String, StringDecoder, StringDecoder](  
    ssc, kafkaParams, Map(KafkaTopicRaw -> 10), StorageLevel.DISK_ONLY_2)  
    .map { case (_, line) => line.split(",") }  
    .map(RawWeatherData(_))  
  
/** Saves the raw data to Cassandra – raw table. */  
stream.saveToCassandra(keyspace, rawDataTable)  
  
stream.map { data =>  
    (data.wsid, data.year, data.month, data.day, data.oneHourPrecip)  
}.saveToCassandra(keyspace, dailyPrecipitationTable)  
  
ssc.start()  
ssc.awaitTermination
```

Streaming Window Operations

```
// where pairs are (word, count)
pairsStream
  .flatMap { case (k,v) => (k,v.value) }
  .reduceByKeyAndWindow((a:Int,b:Int) => (a + b),
    Seconds(30), Seconds(10))
  .saveToCassandra(keyspace,table)
```

Streaming Window Operations

- `window(Duration, Duration)`
- `countByWindow(Duration, Duration)`
- `reduceByWindow(Duration, Duration)`
- `countByValueAndWindow(Duration, Duration)`
- `groupByKeyAndWindow(Duration, Duration)`
- `reduceByKeyAndWindow((V, V) => V, Duration, Duration)`

Ten Things About Cassandra

You always wanted to know but were afraid to ask

DATASTAX

Apache Cassandra

- Elasticity - scale to as many nodes as you need, when you need
- Always On - No single point of failure, Continuous availability
- Masterless peer to peer architecture
- Designed for Replication
- Flexible Data Storage
- Read and write to any node syncs across the cluster
- Operational simplicity - with all nodes in a cluster being the same, there is no complex configuration to manage

Easy to use

- CQL - familiar syntax
- Friendly to programmers
- Paxos for locking

```
CREATE TABLE users (
    username varchar,
    firstname varchar,
    lastname varchar,
    email list<varchar>,
    password varchar,
    created_date timestamp,
    PRIMARY KEY (username)
);
```

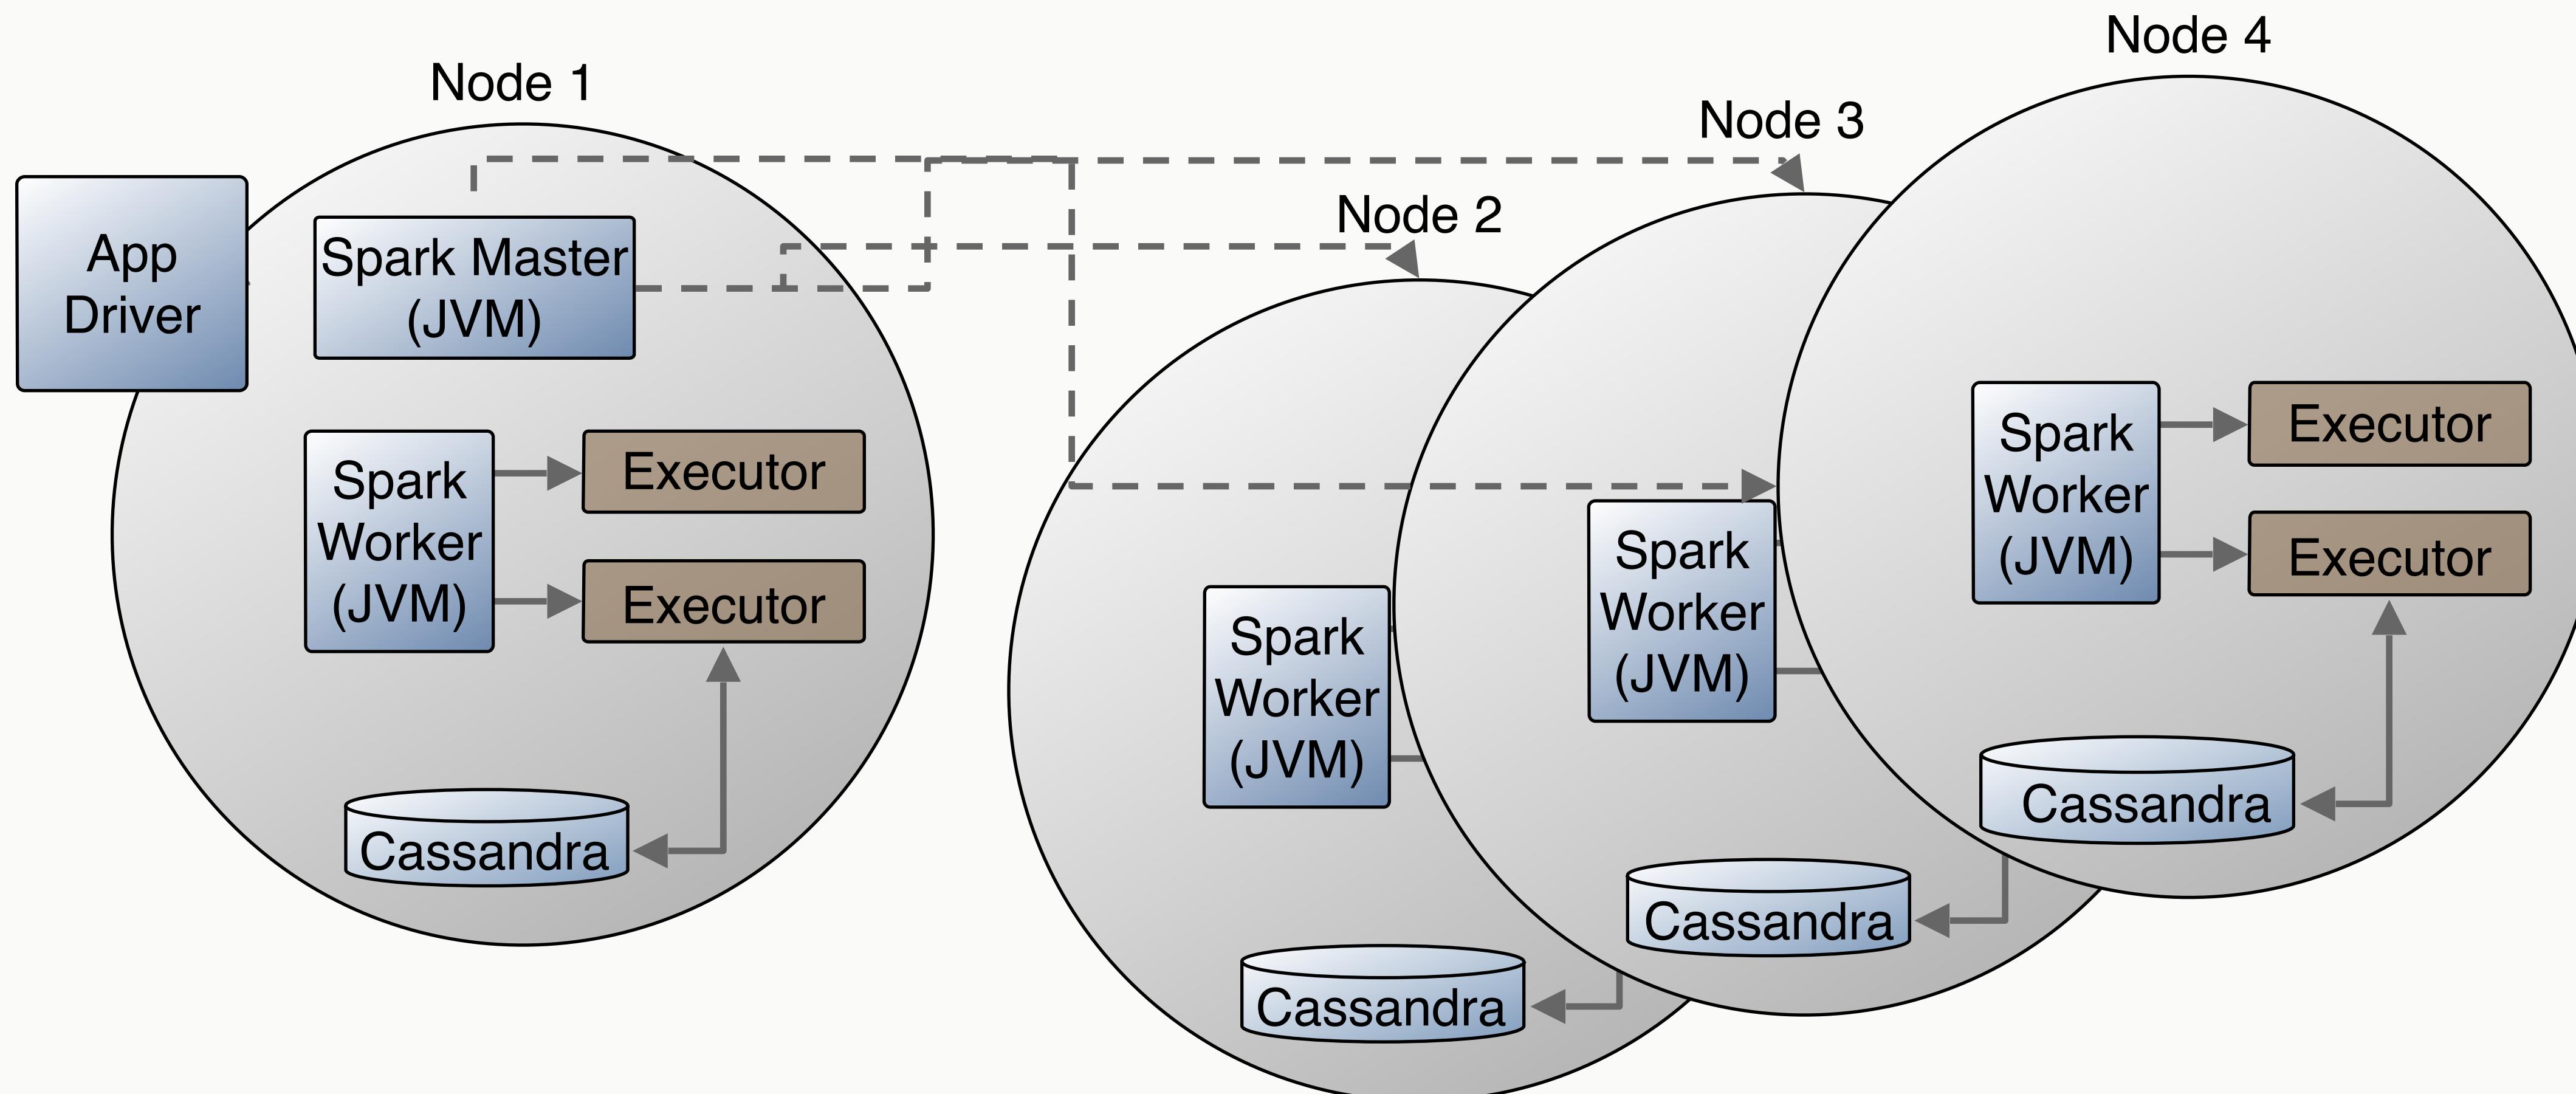
```
INSERT INTO users (username, firstname, lastname,
    email, password, created_date)
VALUES ('hedelson','Helena','Edelson',
    ['helena.edelson@datastax.com'],'ba27e03fd95e507daf2937c937d499ab','2014-11-15 13:50:00');
```

```
INSERT INTO users (username, firstname,
    lastname, email, password, created_date)
VALUES ('pmcfadin','Patrick','McFadin',
    ['patrick@datastax.com'],
    'ba27e03fd95e507daf2937c937d499ab',
    '2011-06-20 13:50:00')
IF NOT EXISTS;
```

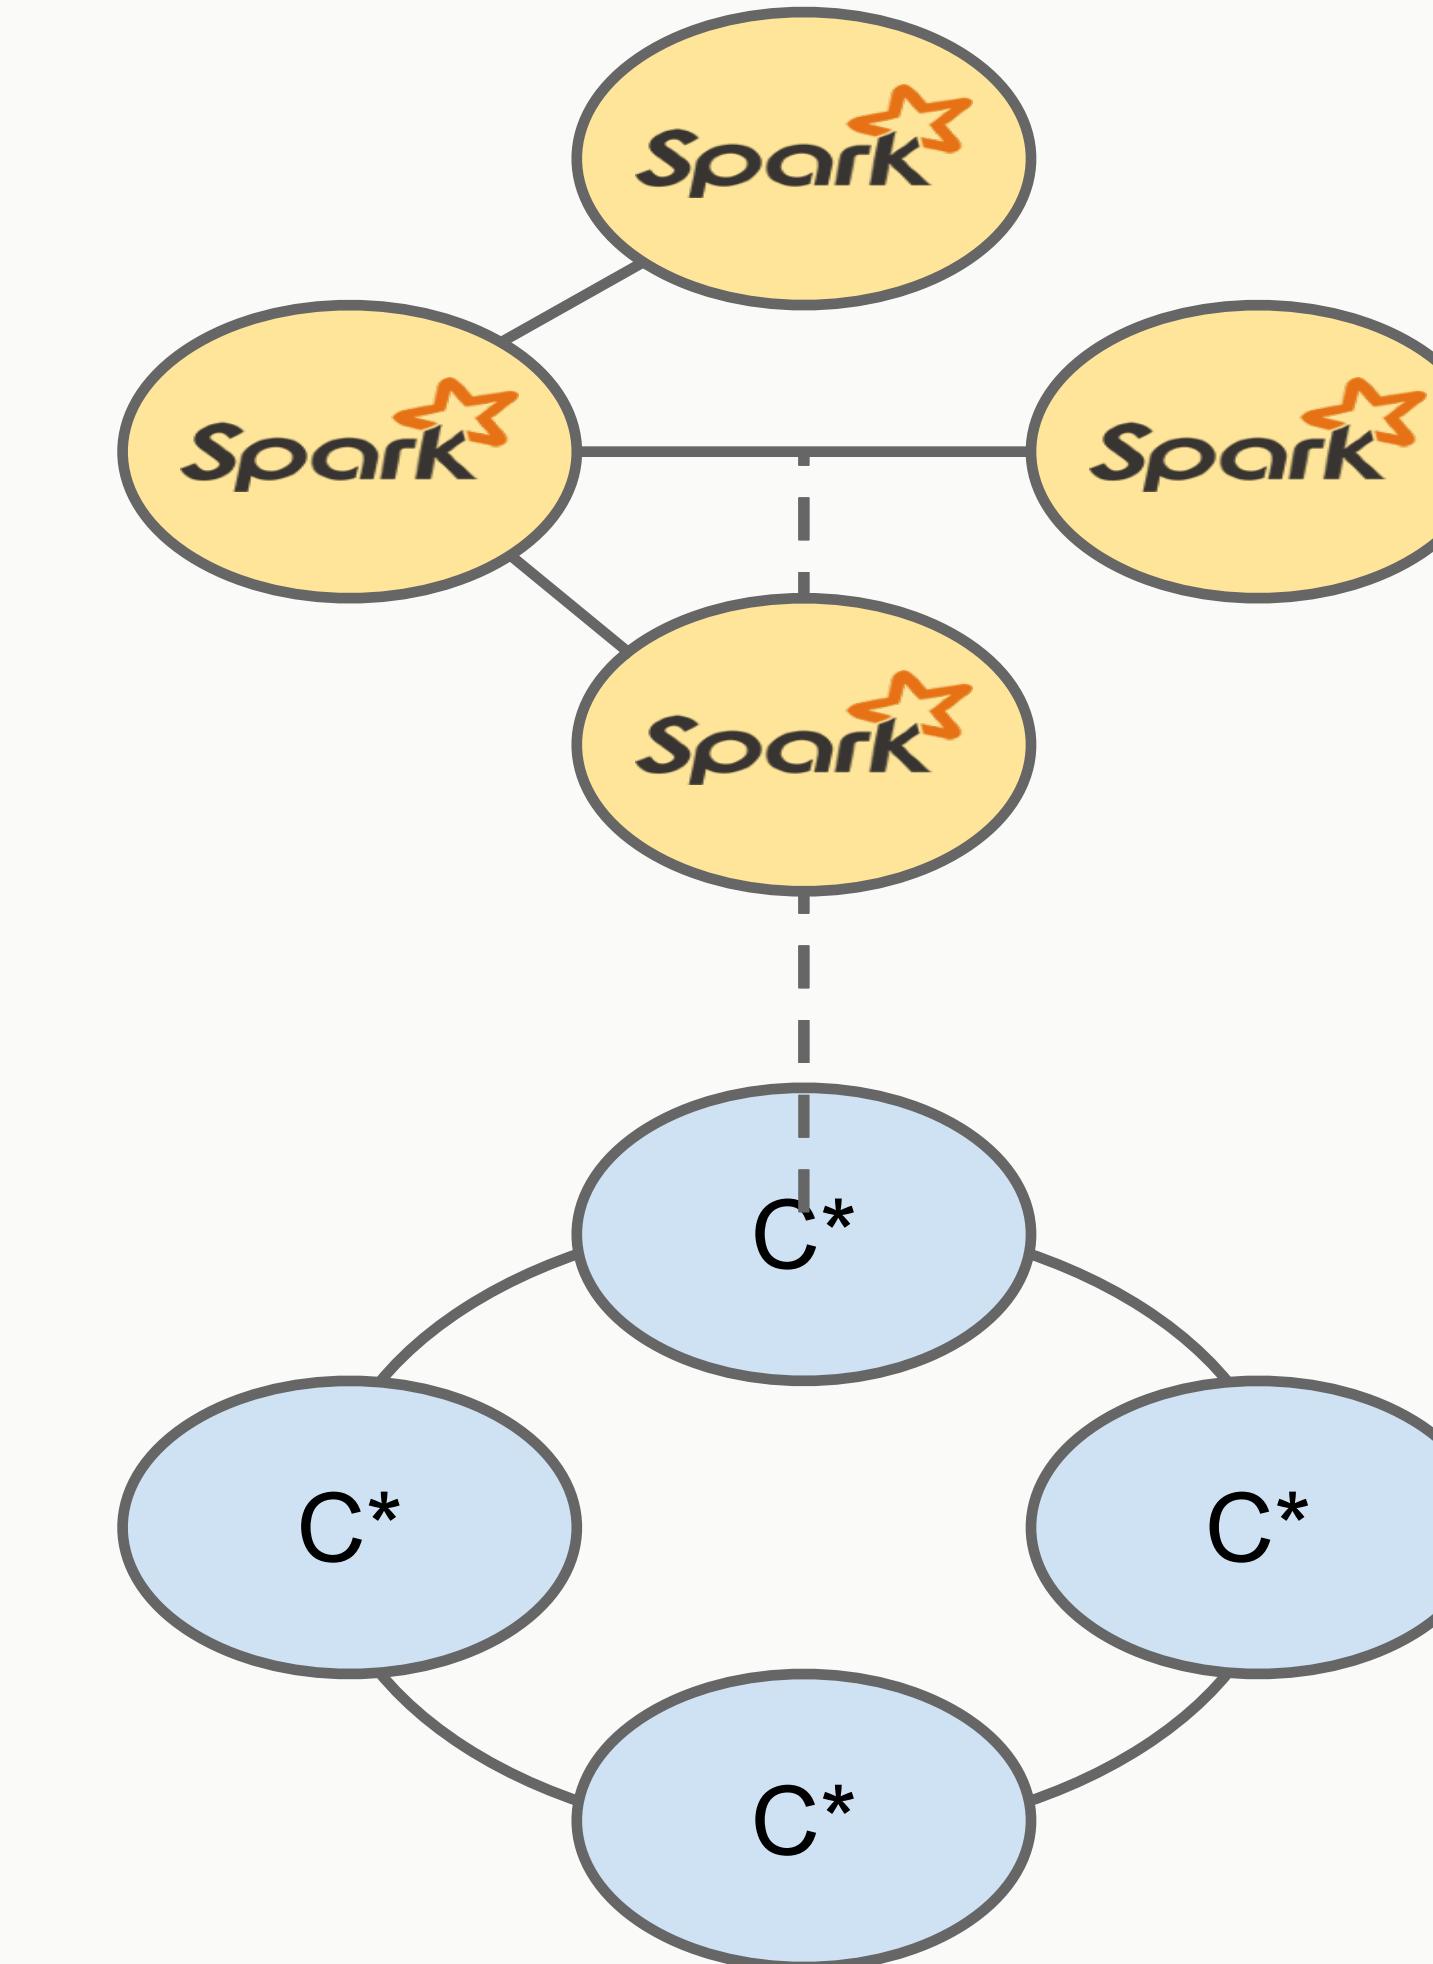
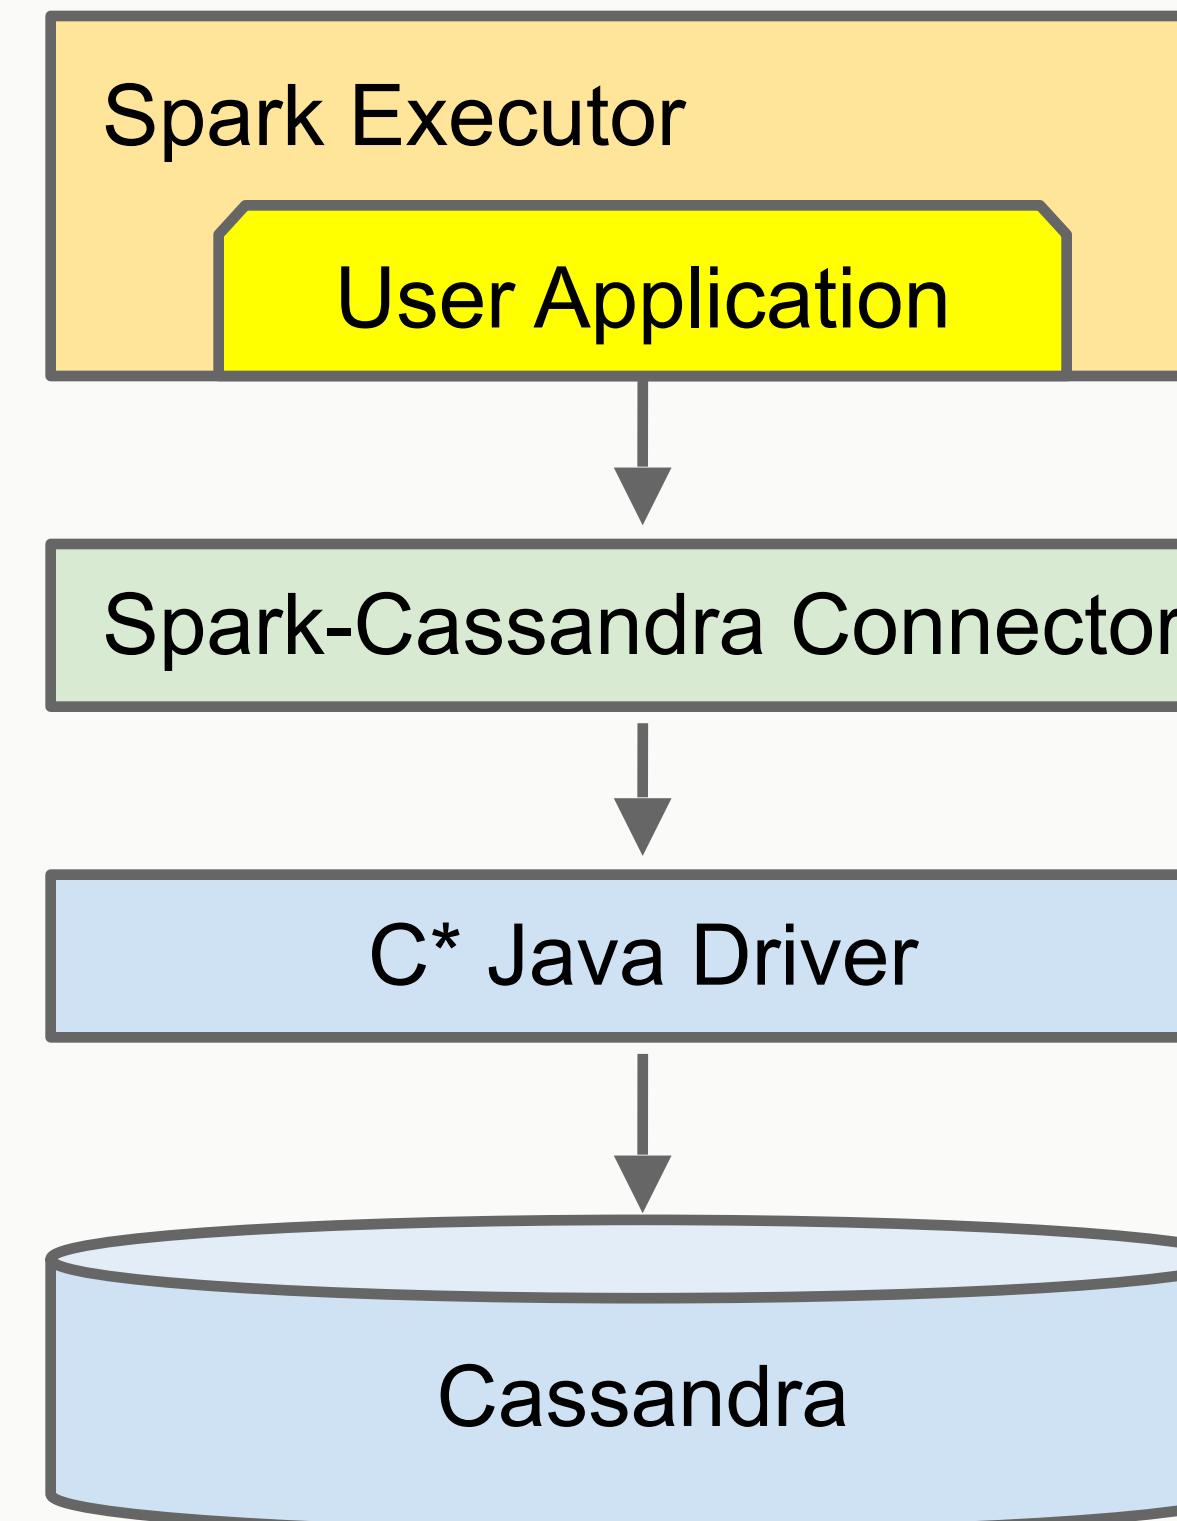
Spark Cassandra Integration



Deployment Architecture



Spark Cassandra Connector

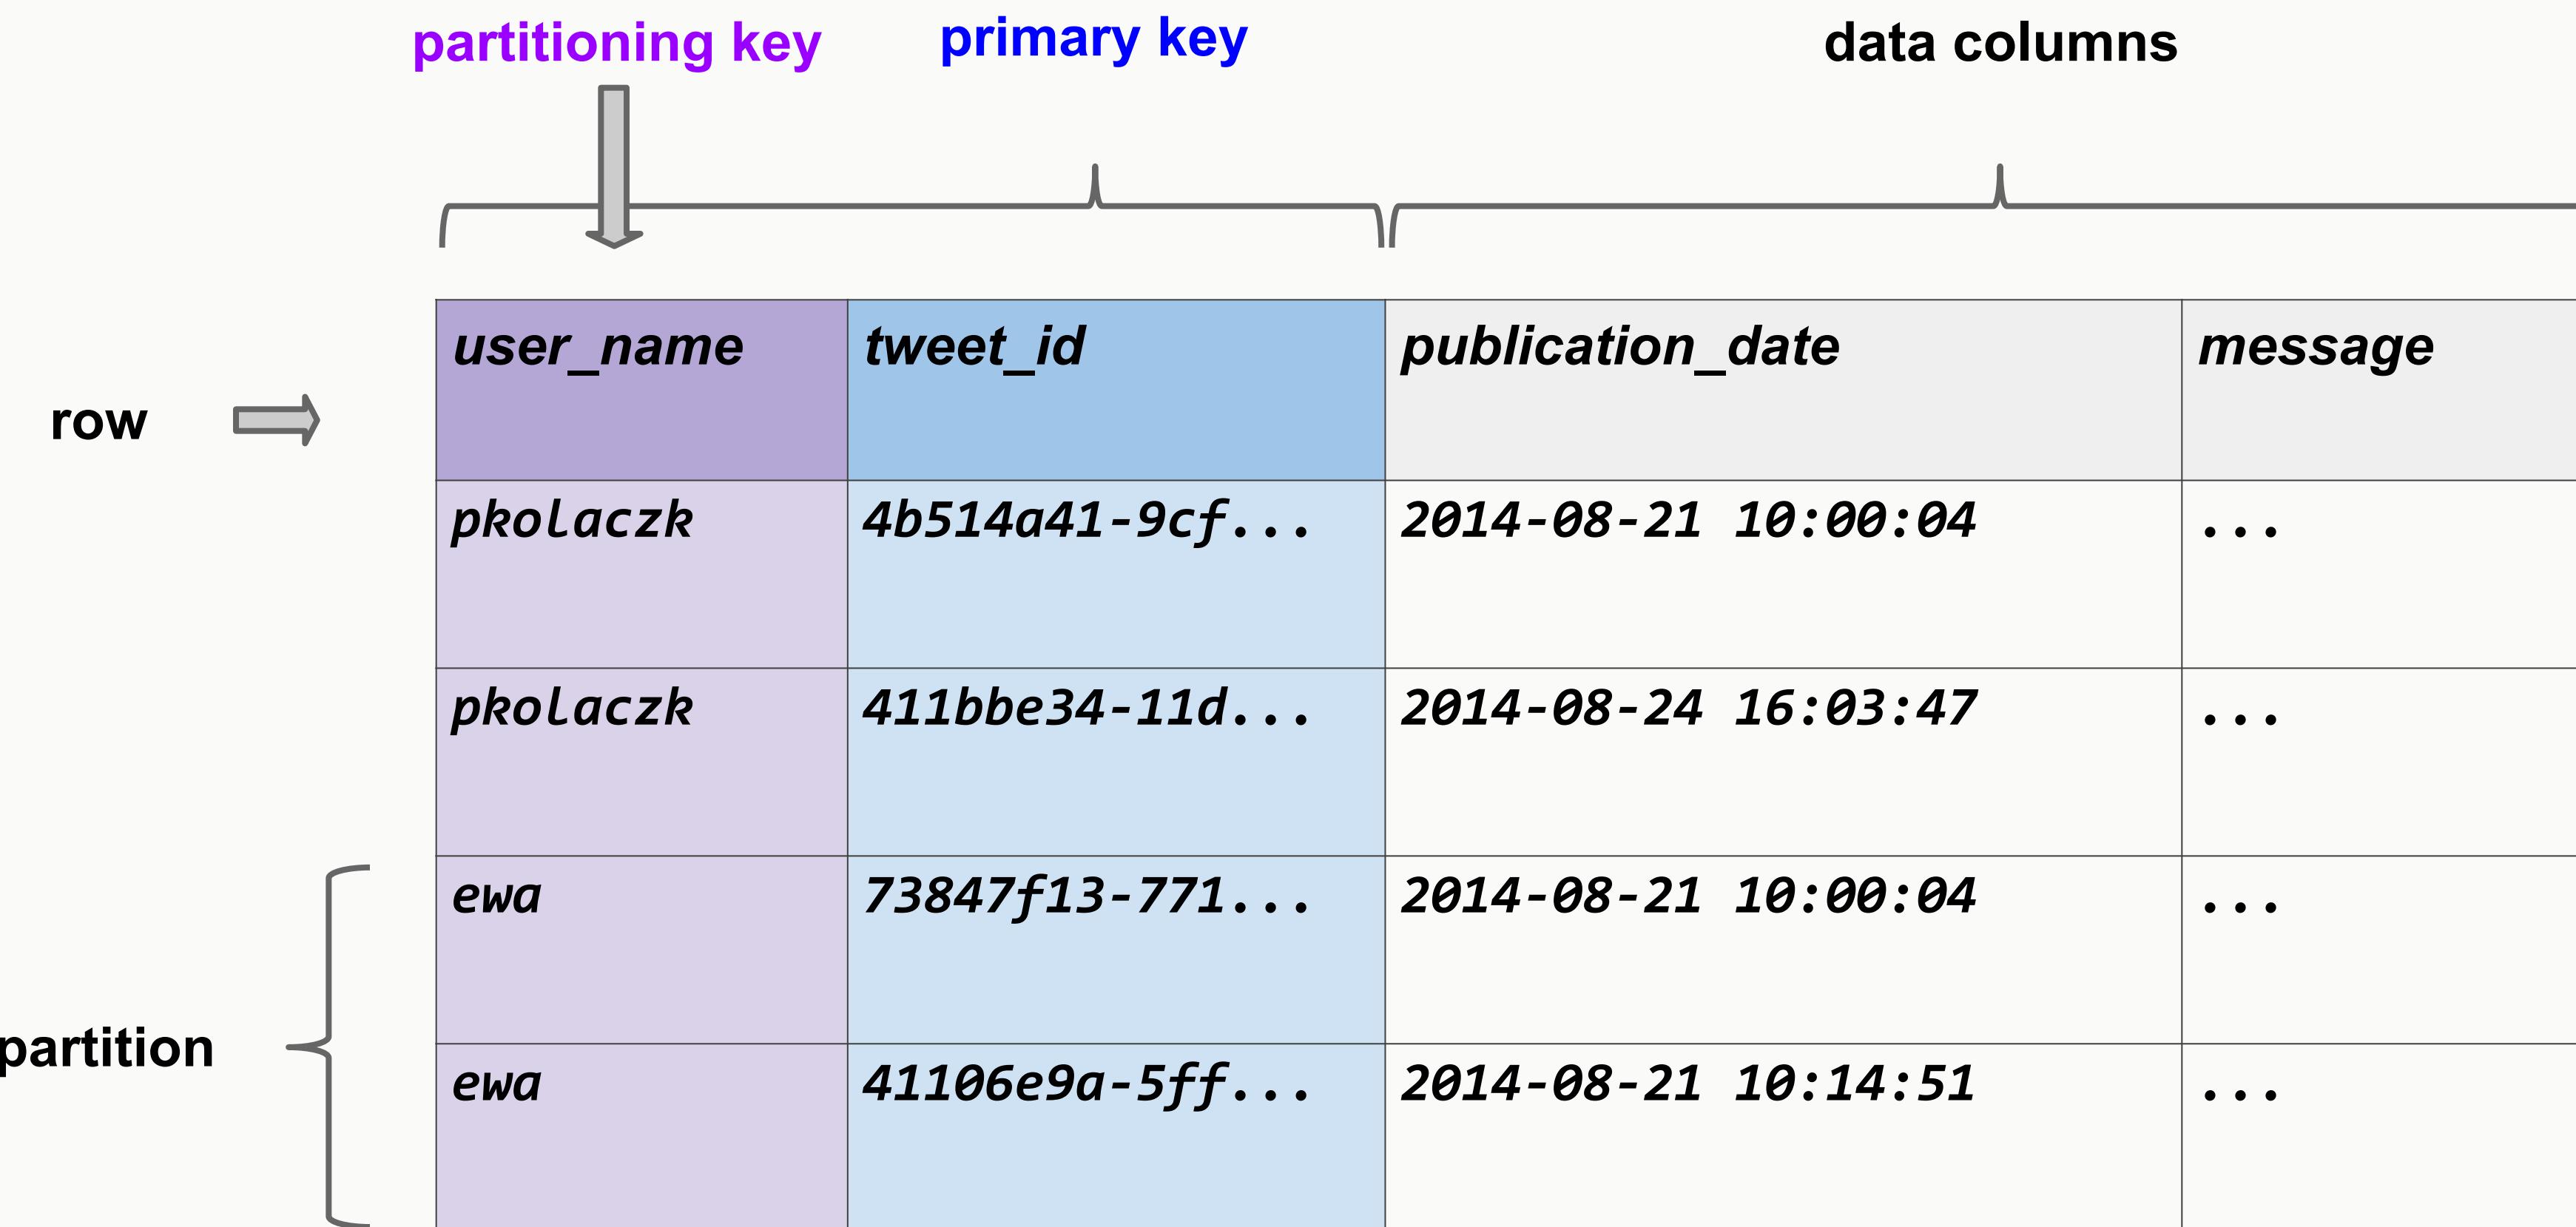


Spark Cassandra Connector

[**https://github.com/datastax/spark-cassandra-connector**](https://github.com/datastax/spark-cassandra-connector)

- Loads data from Cassandra to Spark
 - Writes data from Spark to Cassandra
 - Handles type conversions
 - Offers an object mapper
-
- Implemented in Scala
 - Scala and Java APIs
 - Open Source

Cassandra Data Model



Use Cases

- Get data by weather station
- Get data for a single date and time
- Get data for a range of dates and times
- Compute, store and quickly retrieve daily, monthly and annual aggregations of data

**Design Data Model to support queries (speed, scale, distribution)
vs contort queries to fit data model**

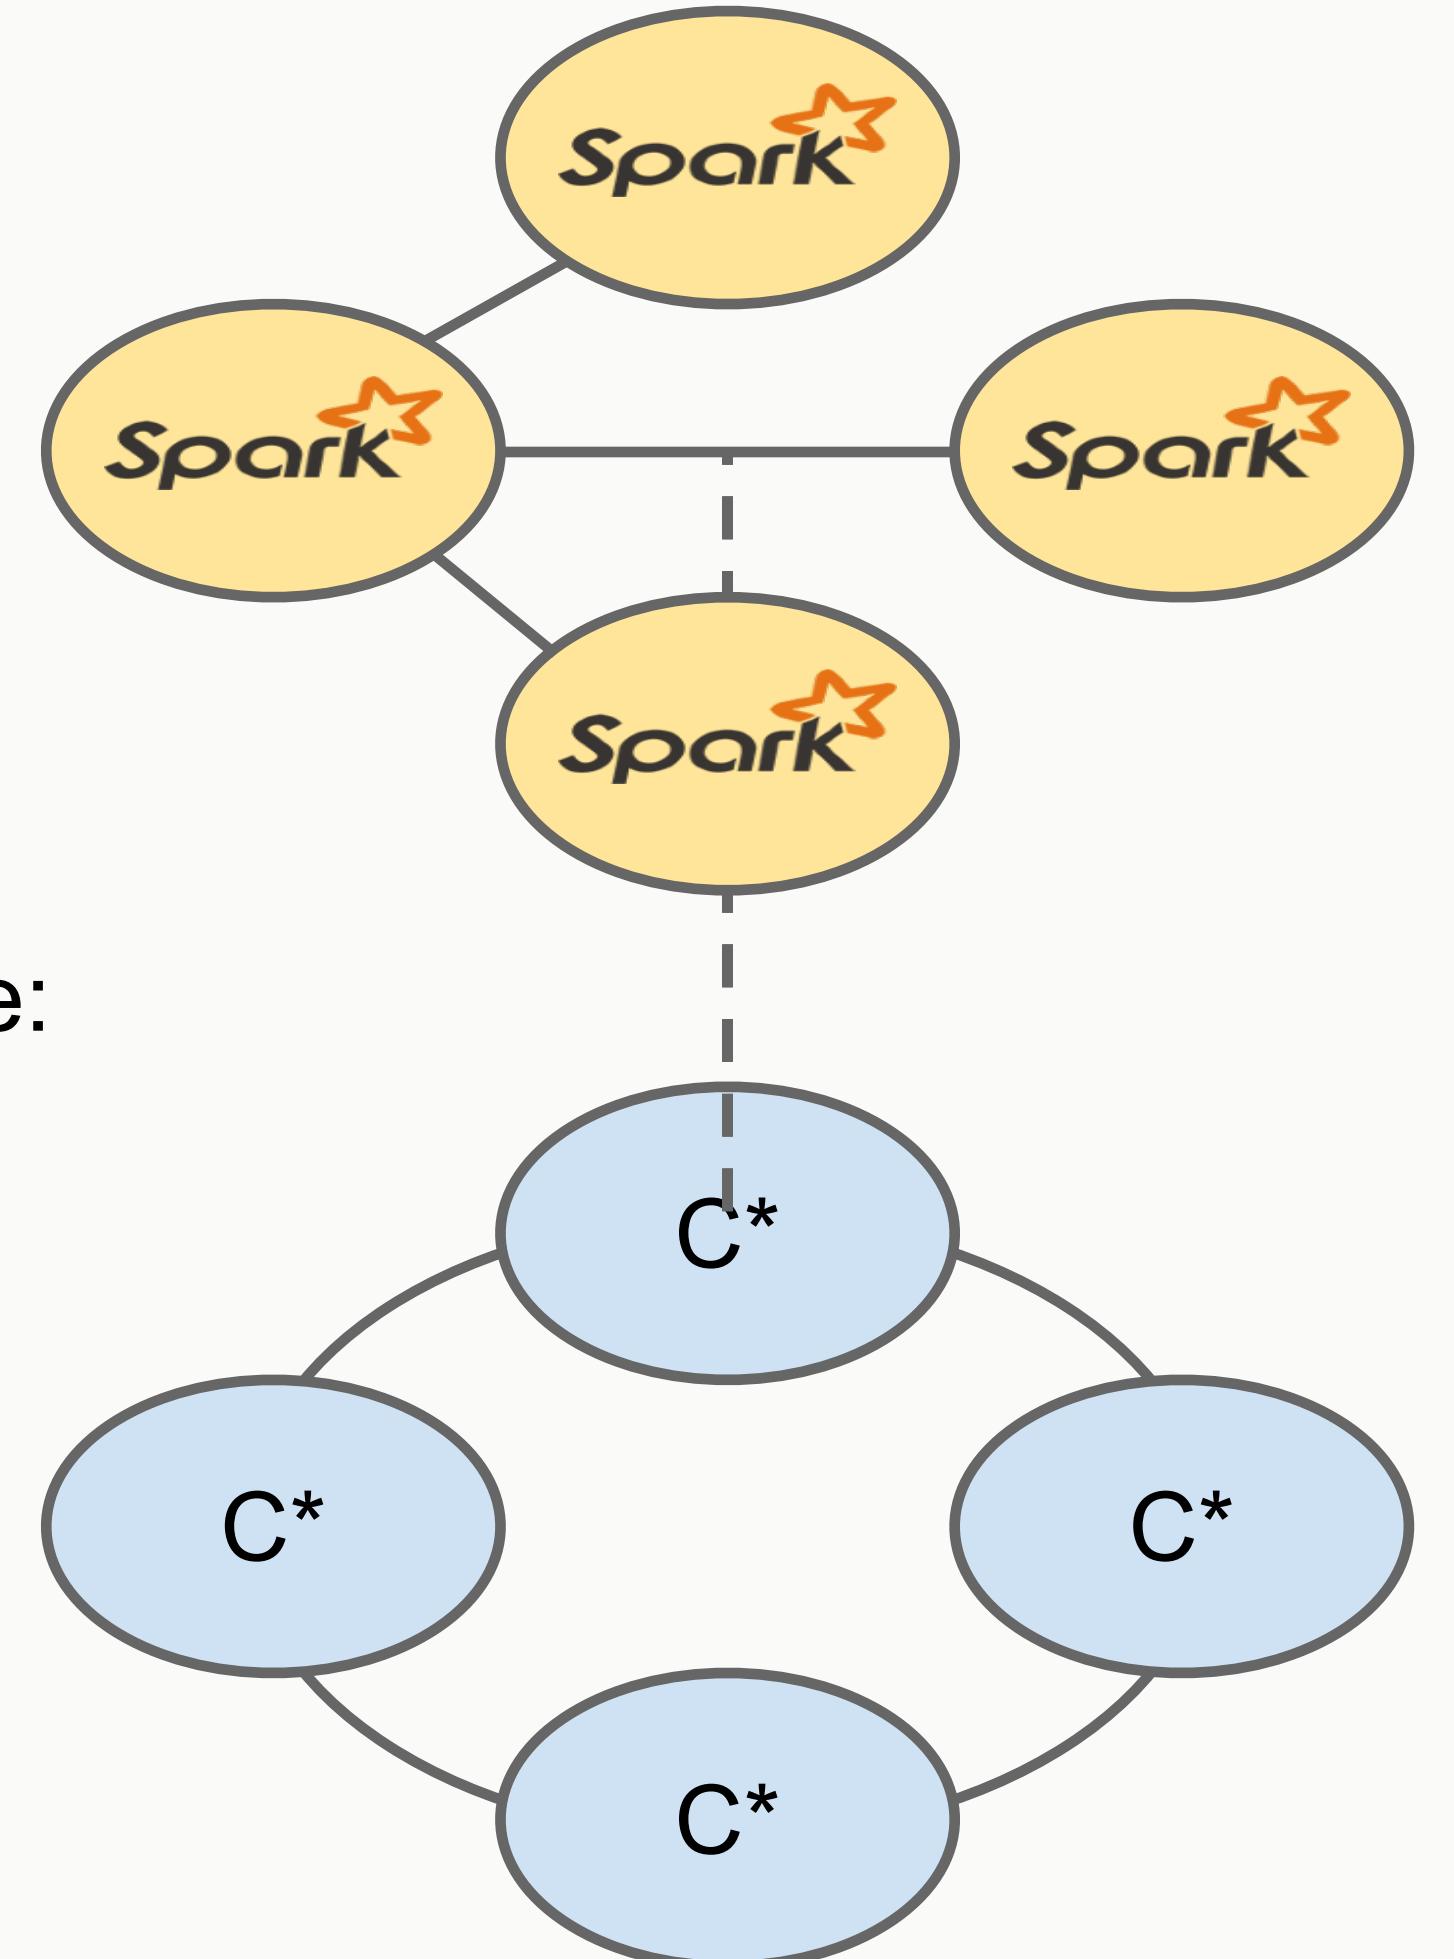
- Store raw data per weather station
- Store time series in order: most recent to oldest
- Compute and store aggregate data in the stream
- Set TTLs on historic data

Data Locality

- Spark asks an RDD for a list of its partitions (splits)
- Each split consists of one or more token-ranges
- For every partition
 - Spark asks RDD for a list of preferred nodes to process on
 - Spark creates a task and sends it to one of the nodes for execution

Every Spark task uses a CQL-like query to fetch data for the given token range:

```
SELECT "key", "value"  
FROM "test"."kv"  
WHERE  
    token("key") > 595597420921139321 AND  
    token("key") <= 595597431194200132  
ALLOW FILTERING
```



Connector Code and Docs



<https://github.com/datastax/spark-cassandra-connector>

The screenshot shows the GitHub repository page for 'spark-cassandra-connector'. The repository has 301 commits, 9 branches, 6 releases, and 10 contributors. It has 159 stars and 37 forks. The master branch is selected. The repository description states: "If you write a Spark application that needs access to Cassandra, this library is for you". The right sidebar shows 35 issues and 2 pull requests.

Add It To Your Project:

```
"com.datastax.spark" %% "spark-cassandra-connector" % "1.1.0"
```

Using Spark With Cassandra



Reading Data: From C* To Spark

row
representation keyspace table

The diagram illustrates the flow of data through a series of four downward arrows. The first arrow points from the text 'server side column and row selection' to the code line '.cassandraTable[CassandraRow]("db", "tweets")'. The second arrow points from the code line '.select("user_name", "message")' to the code line '.where("user_name = ?", "ewa")'. The third arrow points from the code line '.where("user_name = ?", "ewa")' to the word 'table'. The fourth arrow points from the word 'table' to the word 'keyspace'. The word 'row representation' is positioned above the first arrow.

```
val table = sc
    .cassandraTable[CassandraRow]("db", "tweets")
    .select("user_name", "message")
    .where("user_name = ?", "ewa")
```

server side column and row selection {

Spark Streaming, Kafka & Cassandra

```
val conf = new SparkConf(true)
  .setMaster("local[*]")
  .setAppName(getClass.getSimpleName)
  .set("spark.executor.memory", "1g")
  .set("spark.cores.max", "1")
  .set("spark.cassandra.connection.host", "127.0.0.1")
val ssc = new StreamingContext(conf, Seconds(30))
```

} Initialization

```
KafkaUtils.createStream[String, String, StringDecoder, StringDecoder](
  ssc, kafka.kafkaParams, Map(topic -> 1), StorageLevel.MEMORY_ONLY)
  .map(_.value)
  .countByValue()
  .saveToCassandra("mykeyspace", "wordcount")
```

} Transformations
and Action

```
ssc.start()
ssc.awaitTermination()
```

Spark Streaming, Twitter & Cassandra



```
val stream = TwitterUtils.createStream(  
    ssc, auth, filters, StorageLevel.MEMORY_ONLY_SER_2)
```

```
val stream = TwitterUtils.createStream(ssc, auth, Nil,  
StorageLevel.MEMORY_ONLY_SER_2)
```

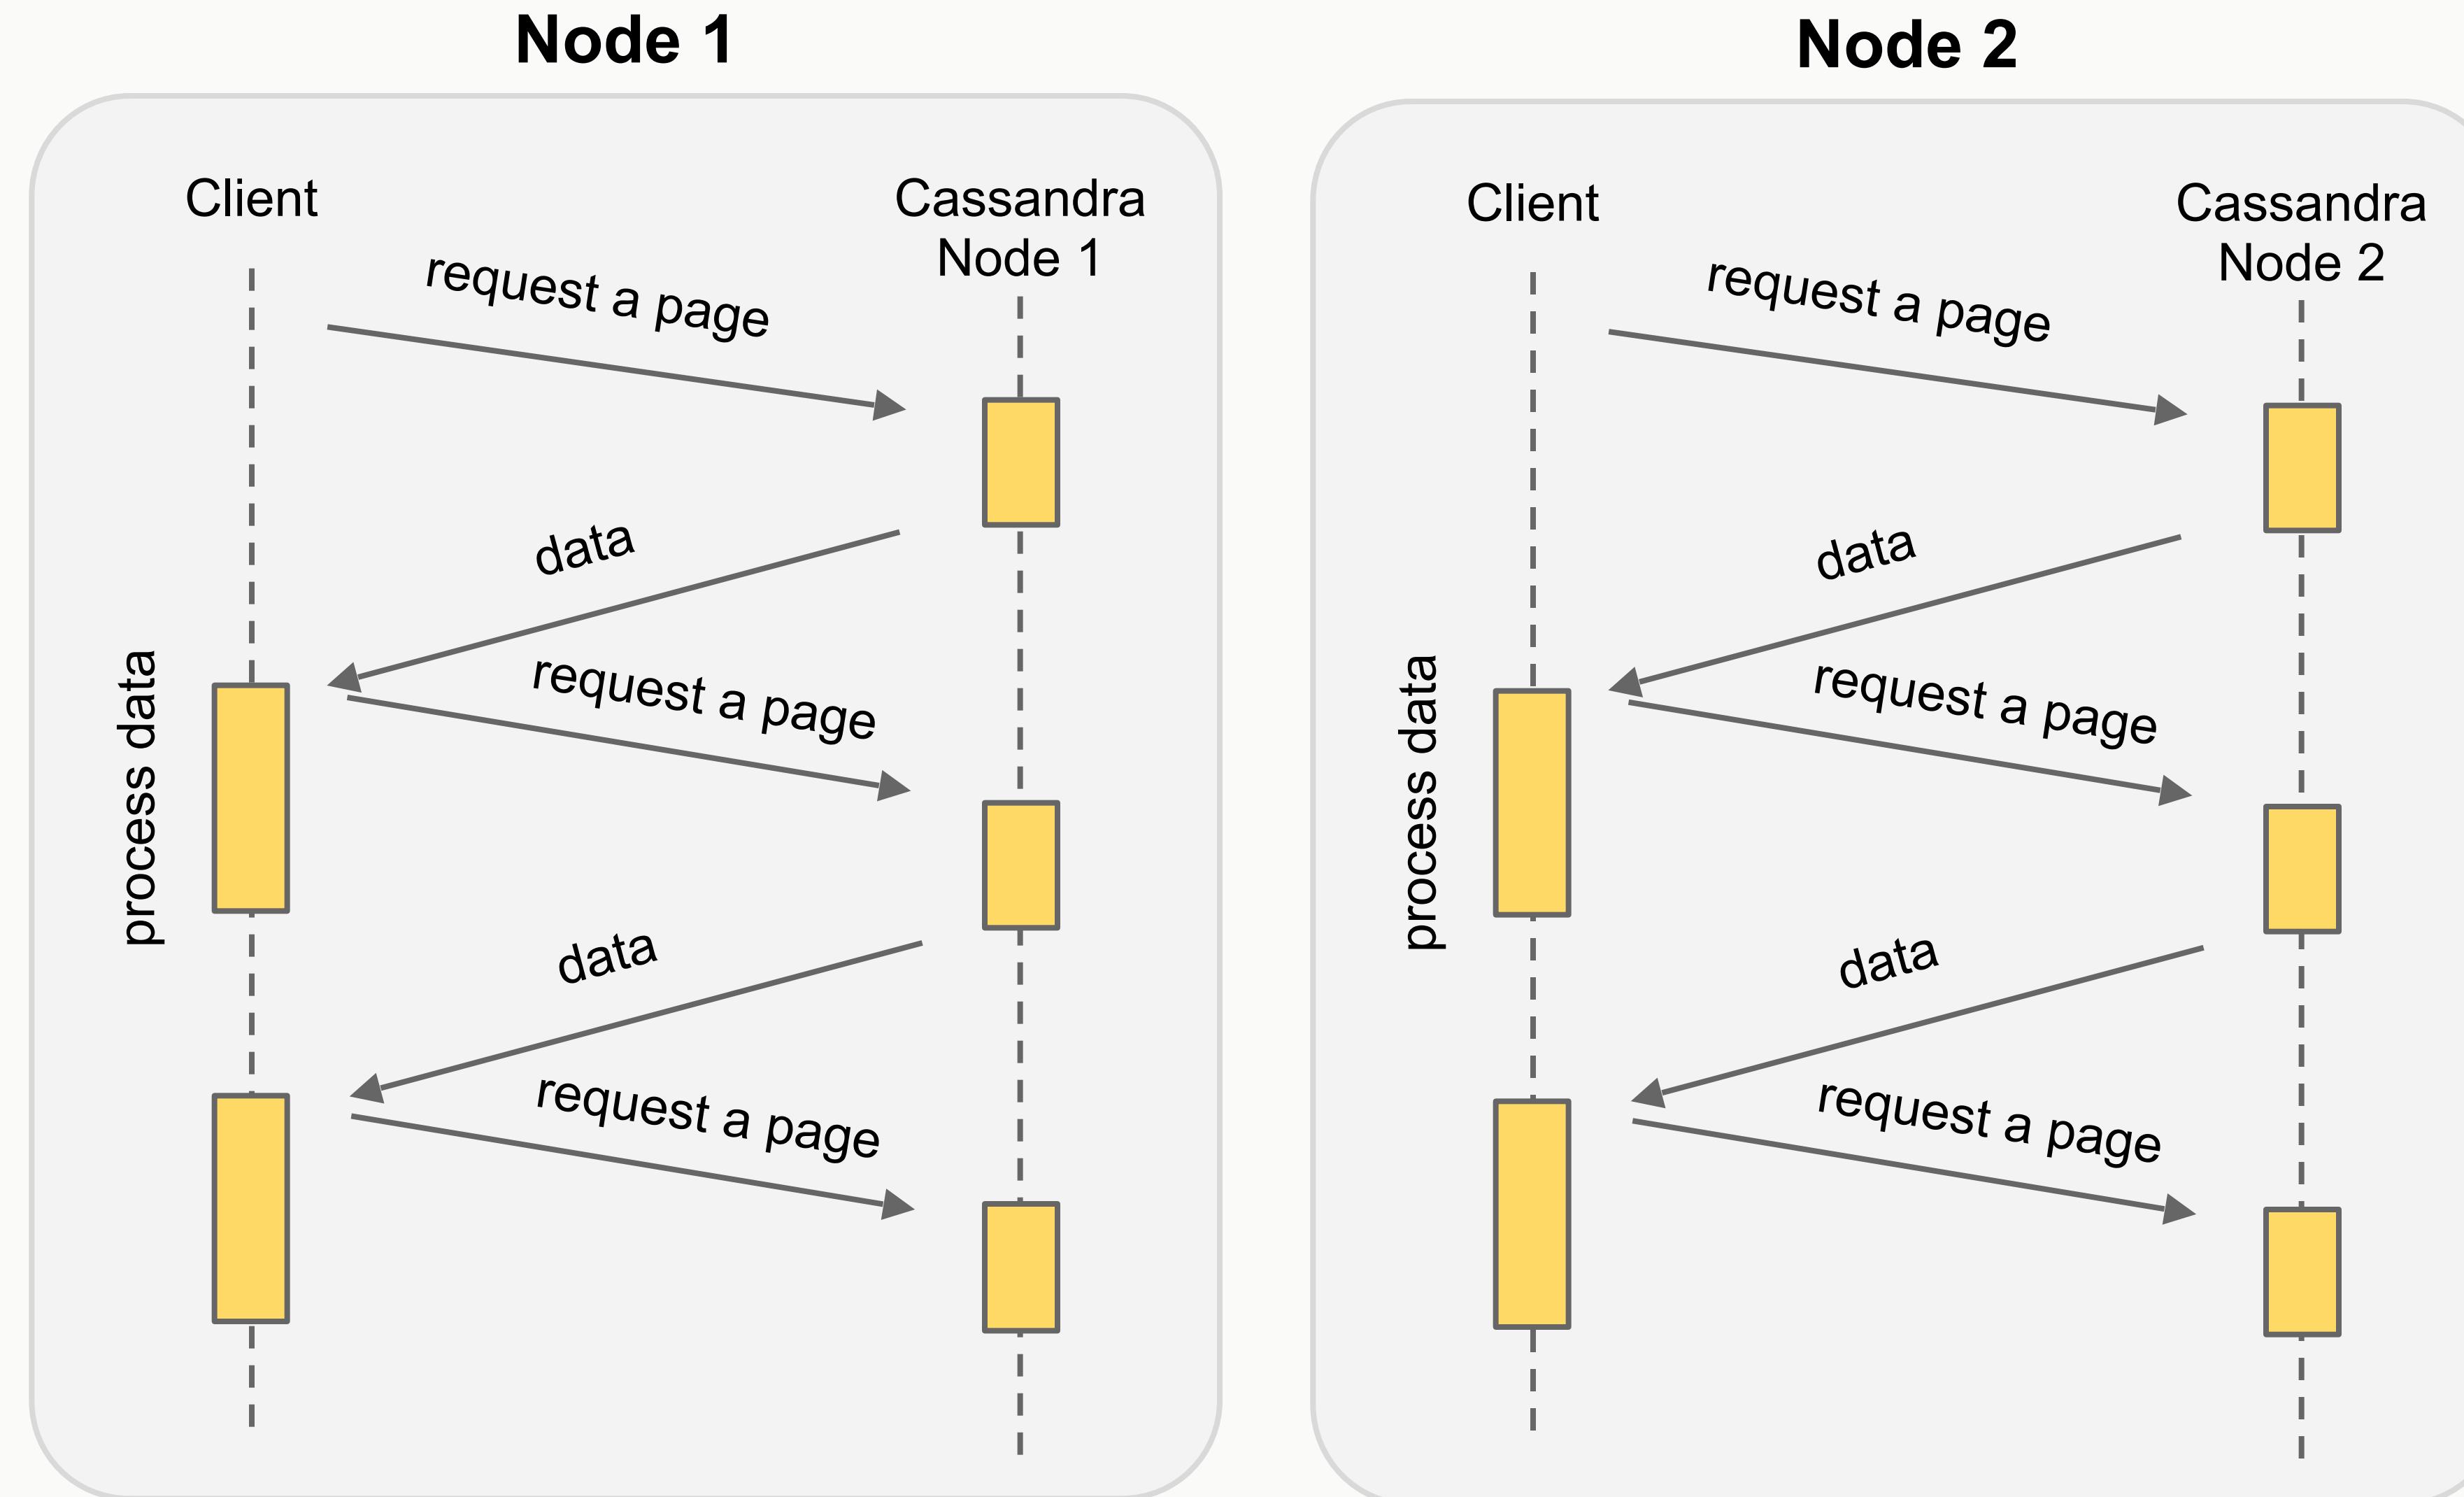
```
/** Note that Cassandra is doing the sorting for you here. */  
stream.flatMap(_.getText.toLowerCase.split(""\s+""))  
.filter(topics.contains(_))  
.countByValueAndWindow(Seconds(5), Seconds(5))  
.transform((rdd, time) =>  
    rdd.map { case (term, count) => (term, count, now(time)) })  
.saveToCassandra(keyspace, table)
```

Paging Reads with .cassandraTable



- Configurable Page Size
- Controls how many CQL rows to fetch at a time, when fetching a single partition
- Connector returns an `Iterator` for rows to Spark
- Spark iterates over this, lazily

ResultSet Paging and Pre-Fetching



Converting Columns to a Type

```
val table = sc.cassandraTable[CassandraRow]("db", "tweets")
val row: CassandraRow = table.first
val userName = row.getString("user_name")
val tweetId = row.getUUID("row_id")
```

Converting Columns To a Type

```
val row: CassandraRow = table.first

val date1 = row.get[java.util.Date]("publication_date")
val date2: org.joda.time.DateTime = row.get[org.joda.time.DateTime]("publication_date")
val date3: org.joda.time.DateTime = row.getDateTime("publication_date")
val date5: Option[DateTime] = row.getDateTimeOption("publication_date")

val list1 = row.get[Set[UUID]]("collection")
val list2 = row.get[List[UUID]]("collection")
val list3 = row.get[Vector[UUID]]("collection")
val list3 = row.get[Vector[String]]("collection")

val nullable = row.get[Option[String]]("nullable_column")
```

Converting Rows To Objects

```
case class Tweet(  
  userName: String,  
  tweetId: UUID,  
  publicationDate: Date,  
  message: String)
```

```
val tweets = sc.cassandraTable[Tweet]("db", "tweets")
```

Scala	Cassandra
message	message
column1	column_1
userName	user_name



Scala	Cassandra
Message	Message
column1	column1
userName	userName

Converting Cassandra Rows To Tuples

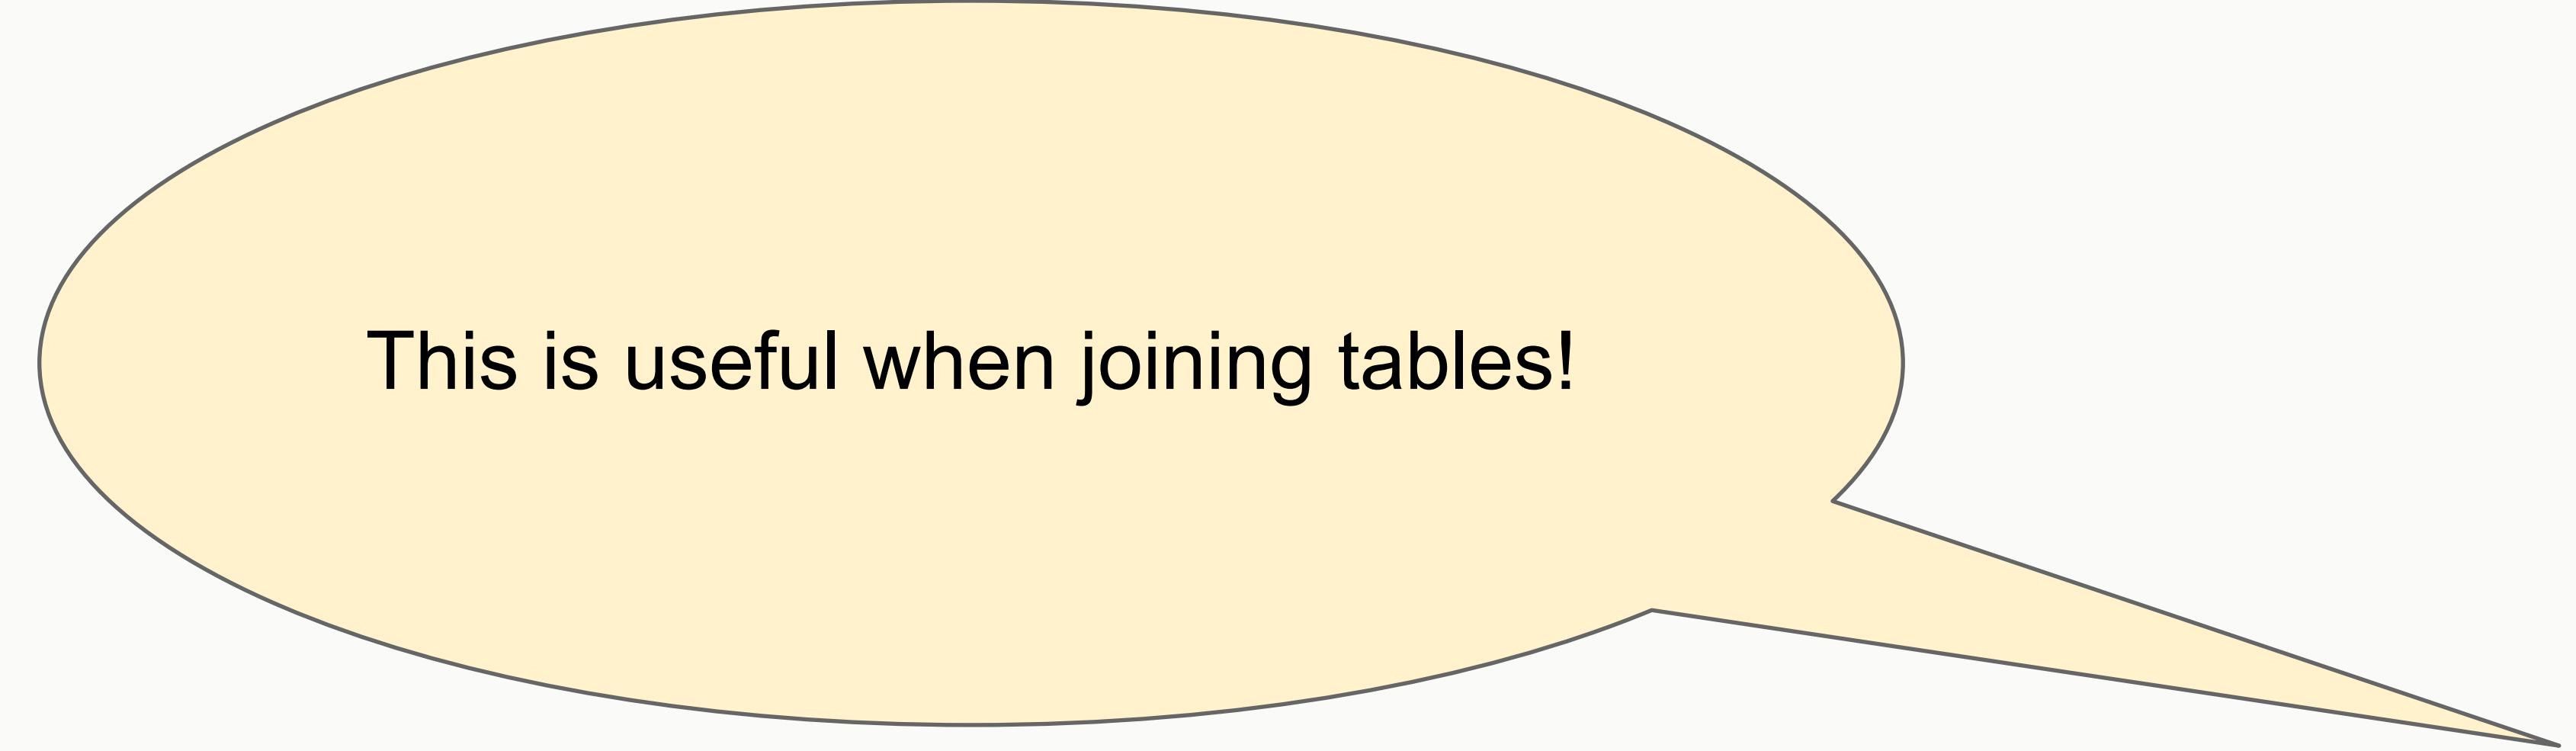
```
val tweets = sc
  .cassandraTable[(Int, Date, String)]("mykeyspace", "clustering")
  .select("cluster_id", "time", "cluster_name")
  .where("time > ? and time < ?", "2014-07-12 20:00:01", "2014-07-12 20:00:03")
```

When returning tuples, always use select to specify the column order.

Converting Rows to Key-Value Pairs

```
case class Key(...)  
case class Value(...)
```

```
val rdd = sc.cassandraTable[(Key, Value)]("keyspace", "table")
```



This is useful when joining tables!

Writing Data

```
cqlsh> CREATE TABLE test.words(word TEXT PRIMARY KEY, count INT);
```

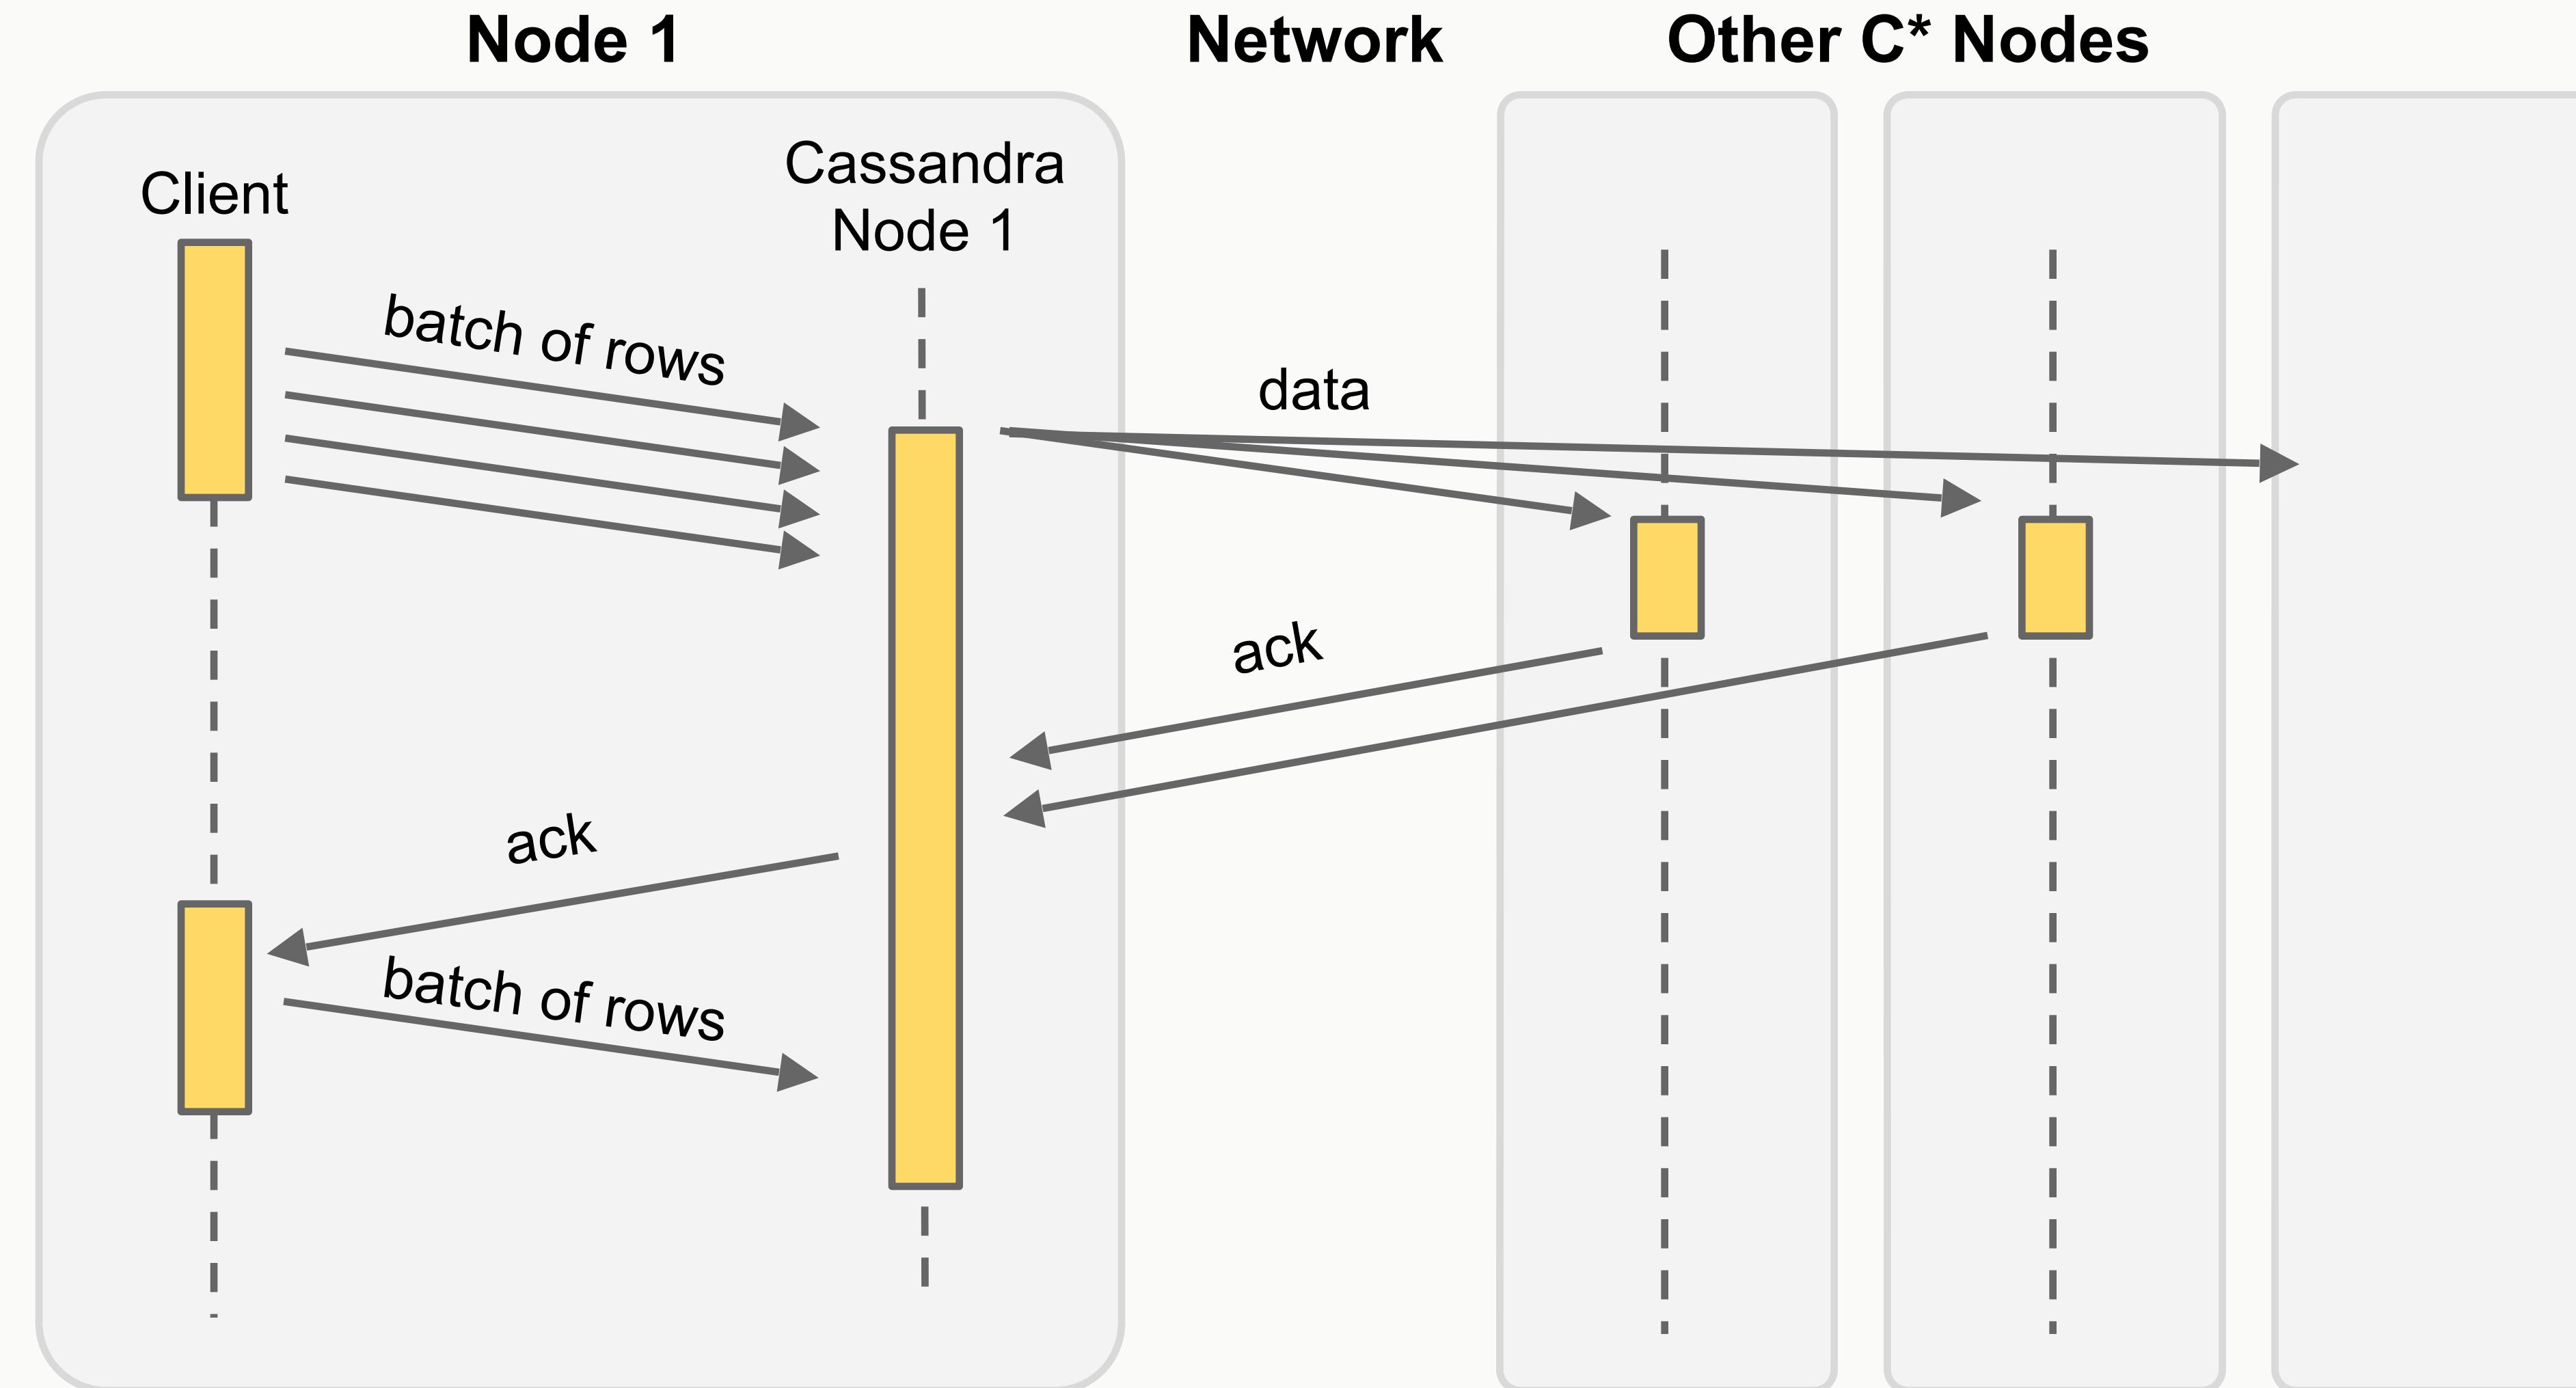
Scala:

```
sc.parallelize(Seq(("foo", 2), ("bar", 5)))
  .saveToCassandra("test", "words", SomeColumns("word", "count"))
```

```
cqlsh:test> select * from words;
```

word	count
bar	5
foo	2

How Writes Are Executed



Other Sweet Features



Spark SQL with Cassandra

NOSQL Joins - Across the spark cluster, to Cassandra (multi-datacenter)

```
val sparkContext = new SparkContext(sparkConf)
val cc = new CassandraSQLContext(sparkContext)
cc.setKeyspace("nosql_joins")
```

```
cc.sql("""
    SELECT test1.a, test1.b, test1.c, test2.a
    FROM test1 AS test1
    JOIN
        test2 AS test2 ON test1.a = test2.a
        AND test1.b = test2.b
        AND test1.c = test2.c
    """).map(Data(_))
    .saveToCassandra("nosql_joins", "table3")
```

Spark SQL: Txt, Parquet, JSON Support

```
import com.datastax.spark.connector._  
import org.apache.spark.sql.{Row, SQLContext}  
  
val sql = new SQLContext(sc)  
val json = sc.parallelize(Seq(  
    """{"user":"helena","commits":98, "month":12, "year":2014}""",  
    """{"user":"pkolaczk", "commits":42, "month":12, "year":2014}""")  
  
sql.jsonRDD(json).map(MonthlyCommits(_)).flatMap(doSparkWork())  
    .saveToCassandra("githubstats", "monthly_commits")  
  
sc.cassandraTable[MonthlyCommits]("githubstats", "monthly_commits")  
    .collect foreach println  
  
sc.stop()
```

What's New In Spark?

- Petabyte sort record
 - Myth busting:
 - Spark is in-memory. It doesn't work with big data
 - It's too expensive to buy a cluster with enough memory to fit our data
- Application Integration
 - Tableau, Trifacta, Talend, ElasticSearch, Cassandra
- Ongoing development for Spark 1.2
 - Python streaming, new MLlib API, Yarn scaling...



Recent / Pending Spark Updates

- Usability, Stability & Performance Improvements
 - Improved Lightning Fast Shuffle!
 - Monitoring the performance long running or complex jobs
- Public types API to allow users to create SchemaRDD's
- Support for registering Python, Scala, and Java lambda functions as UDF
- Dynamic bytecode generation significantly speeding up execution for queries that perform complex expression evaluation

Recent Spark Streaming Updates

- Apache Flume: a new pull-based mode (simplifying deployment and providing high availability)
- The first of a set of streaming machine learning algorithms is introduced with streaming linear regression.
- Rate limiting has been added for streaming inputs

Recent MLlib Updates

- New library of statistical packages which provides exploratory analytic functions *stratified sampling, correlations, chi-squared tests, creating random datasets...)
- Utilities for feature extraction (Word2Vec and TF-IDF) and feature transformation (normalization and standard scaling).
- Support for nonnegative matrix factorization and SVD via Lanczos.
- Decision tree algorithm has been added in Python and Java.
- Tree aggregation primitive
- Performance improves across the board, with improvements of around

Recent/Pending Connector Updates



- Spark SQL (came in 1.1)
- Python API
- Official Scala Driver for Cassandra
- Removing last traces of Thrift!!
- Performance Improvements
 - Token-aware data repartitioning
 - Token-aware saving
 - Wide-row support - no costly groupBy call

Resources

- <http://spark.apache.org>
- <https://github.com/datastax/spark-cassandra-connector>
- <http://cassandra.apache.org>
- <http://kafka.apache.org>
- <http://akka.io>
- <https://groups.google.com/a/lists.datastax.com/forum/#!forum/spark-connector-user>

www.spark-summit.org



New York March 18-19, 2015



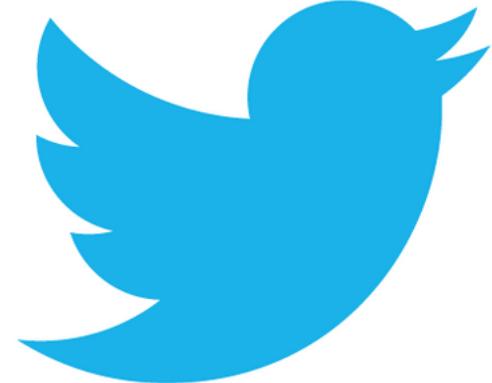
San Francisco June 15-17, 2015

Thanks for listening!



SEPTEMBER 10 - 11, 2014 | SAN FRANCISCO, CALIF. | THE WESTIN ST. FRANCIS HOTEL

Cassandra Summit



@helenaedelson

- Follow-Up Blog Post
- slideshare.net/helenaedelson

DATASTAX

The DataStax logo consists of the word "DATASTAX" in a bold, black, sans-serif font. To the right of the "X", there is a graphic element composed of four teal-colored circles of increasing size, arranged in a curve that tapers towards the right.